

# 20104016

## DEENA

### Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2006.csv")
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2006-02-01 01:00:00	NaN	1.84	NaN	NaN	NaN	155.100006	490.100006	NaN	4.880000	97.
1	2006-02-01 01:00:00	1.68	1.01	2.38	6.36	0.32	94.339996	229.699997	3.04	7.100000	25.
2	2006-02-01 01:00:00	NaN	1.25	NaN	NaN	NaN	66.800003	192.000000	NaN	4.430000	34.
3	2006-02-01 01:00:00	NaN	1.68	NaN	NaN	NaN	103.000000	407.799988	NaN	4.830000	28.
4	2006-02-01 01:00:00	NaN	1.31	NaN	NaN	NaN	105.400002	269.200012	NaN	6.990000	54.
...	...	...	...	...	...	...	...	...	...	...	...
230563	2006-05-01 00:00:00	5.88	0.83	6.23	NaN	0.20	112.500000	218.000000	NaN	24.389999	93.
230564	2006-05-01 00:00:00	0.76	0.32	0.48	1.09	0.08	51.900002	54.820000	0.61	48.410000	29.
230565	2006-05-01 00:00:00	0.96	NaN	0.69	NaN	0.19	135.100006	179.199997	NaN	11.460000	64.
230566	2006-05-01 00:00:00	0.50	NaN	0.67	NaN	0.10	82.599998	105.599998	NaN	NaN	94.
230567	2006-05-01 00:00:00	1.95	0.74	1.99	4.00	0.24	107.300003	160.199997	2.01	17.730000	52.

230568 rows × 17 columns

# Data Cleaning and Data Preprocessing

In [3]: `df = df.dropna()`

In [4]: `df.columns`

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3',  
'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'],  
dtype='object')

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   date        24758 non-null  object 
1   BEN         24758 non-null  float64
2   CO          24758 non-null  float64
3   EBE         24758 non-null  float64
4   MXY         24758 non-null  float64
5   NMHC        24758 non-null  float64
6   NO_2        24758 non-null  float64
7   NOx         24758 non-null  float64
8   OXY         24758 non-null  float64
9   O_3         24758 non-null  float64
10  PM10        24758 non-null  float64
11  PM25        24758 non-null  float64
12  PXY         24758 non-null  float64
13  SO_2        24758 non-null  float64
14  TCH         24758 non-null  float64
15  TOL         24758 non-null  float64
16  station     24758 non-null  int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

```
In [6]: data=df[['CO' , 'station']]
```

```
Out[6]:
```


	CO	station
5	1.69	28079006
22	0.79	28079024
25	1.47	28079099
31	0.85	28079006
48	0.79	28079024
...	...	...
230538	0.40	28079024
230541	0.94	28079099
230547	1.06	28079006
230564	0.32	28079024
230567	0.74	28079099

24758 rows × 2 columns

## Line chart

```
In [7]: data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:~>, <AxesSubplot:~>], dtype=object)
```



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```




## Bar chart

```
In [9]: k = data[0:50]
```

```
In [10]: k.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



## Box chart

In [13]: `data.plot_box()`


Out[13]: <AxesSubplot:>



## Pie chart

In [14]: `plot_stations(station)`


Out[14]: `<AxesSubplot:ylabel='station'>`



**Scatter chart**

In [15]: `data.plot.scatter(x='CO', y='station')`

Out[15]: `<AxesSubplot:xlabel='CO', ylabel='station'>`



In [16]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        24758 non-null  object
1   BEN         24758 non-null  float64
2   CO          24758 non-null  float64
3   EBE         24758 non-null  float64
4   MXY         24758 non-null  float64
5   NMHC        24758 non-null  float64
6   NO_2        24758 non-null  float64
7   NOx         24758 non-null  float64
8   OXY         24758 non-null  float64
9   O_3         24758 non-null  float64
10  PM10        24758 non-null  float64
11  PM25        24758 non-null  float64
12  PXY         24758 non-null  float64
13  SO_2        24758 non-null  float64
14  TCU         24758 non-null  float64
```



In [17]: `df.describe()`

Out[17]:

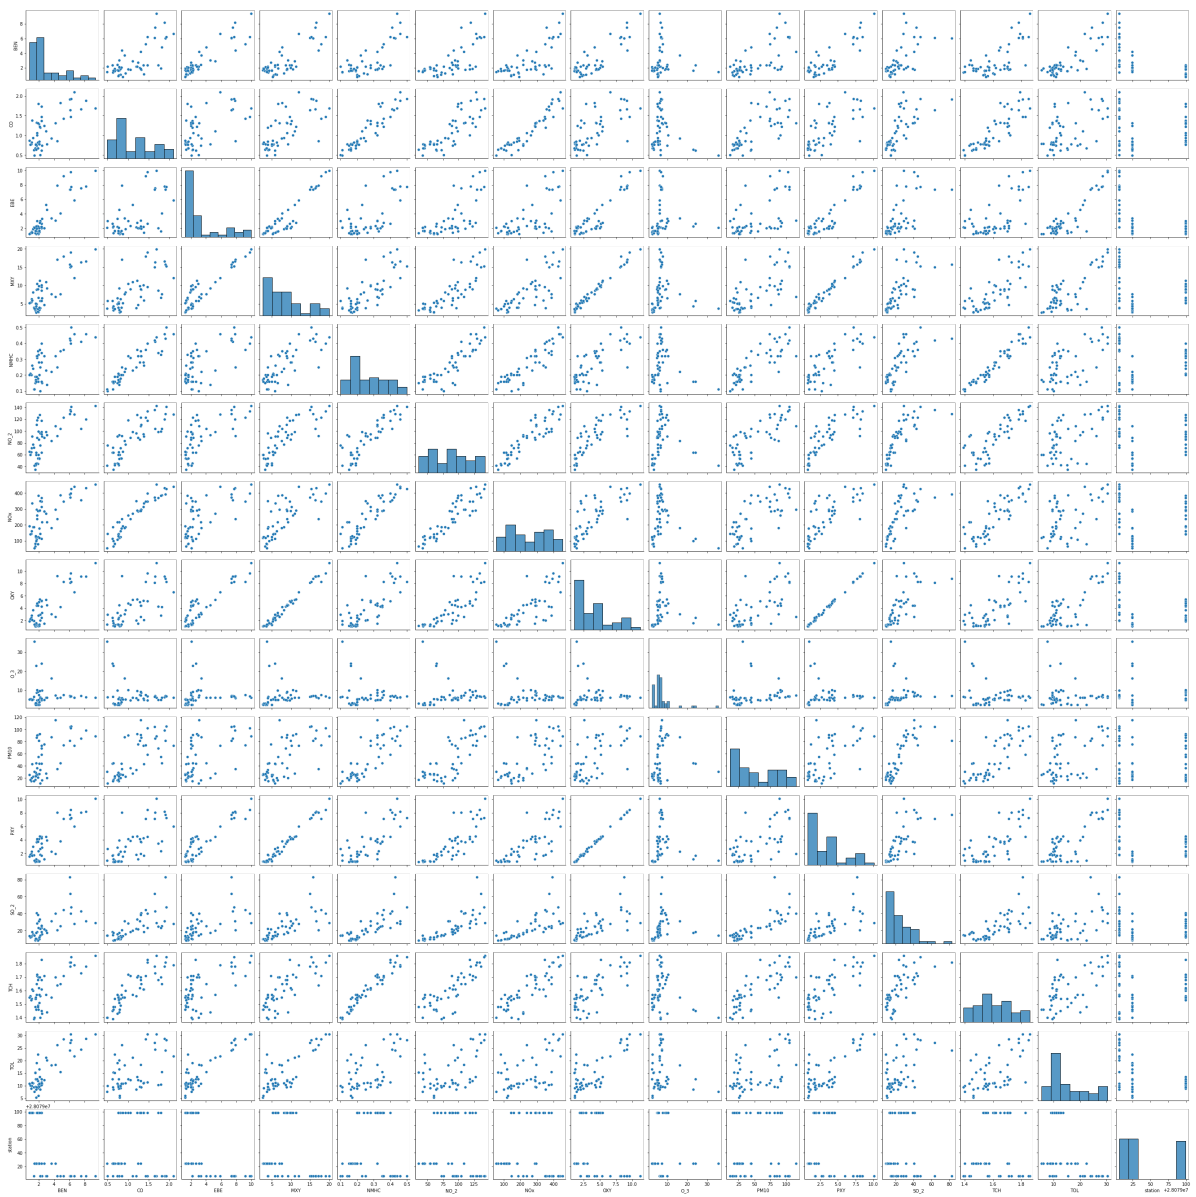
	BEN	CO	EBE	MXY	NMHC	NO_2	
count	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	247
mean	1.350624	0.600713	1.824534	3.835034	0.176546	58.333481	1
std	1.541636	0.419048	1.868939	4.069036	0.126683	40.529382	1
min	0.110000	0.000000	0.170000	0.150000	0.000000	1.680000	
25%	0.450000	0.360000	0.810000	1.060000	0.100000	28.450001	
50%	0.850000	0.500000	1.130000	2.500000	0.150000	52.959999	
75%	1.680000	0.720000	2.160000	5.090000	0.220000	79.347498	1
max	45.430000	7.250000	57.799999	66.900002	2.020000	461.299988	16

In [18]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PM2.5', 'SO_2', 'TCH', 'TOL', 'Station']]`

## EDA AND VISUALIZATION

In [19]: `sns.pairplot(df1[0:50])`

Out[19]: `<seaborn.axisgrid.PairGrid at 0x20335382b80>`



```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
In [23]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train, y_train)
```

Out[24]: LinearRegression()

```
In [25]: lr.intercept
```

Out[25]: 28079013.414891884


```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[26]:

	Co-efficient
BEN	-18.561713
CO	-13.066381
EBE	-21.247476
MXV	2.635001
NMHC	130.910680
NO_2	-0.008165
NOx	-0.009275
OXY	18.118339
O_3	-0.056919
PM10	0.143264
PXY	6.165699
SO_2	-0.653803
TCH	23.167821
TOL	-0.699416

```
In [27]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x20343cfab80>
```



## ACCURACY

```
In [28]: lr.score(y_test, y_test)
```

```
Out[28]: 0.4082591915094467
```

```
In [29]: lr.score(y_train, y_train)
```

```
Out[29]: 0.38629705567566075
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge, Lasso
```

```
In [31]: rr=Ridge(alpha=10)
rr.fit(y_train, y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(y_test, y_test)
```

```
Out[32]: 0.40819618490318543
```

```
In [33]: rr.score(y_train, y_train)
```

```
Out[33]: 0.3856182327593829
```

```
In [34]: la=Lasso(alpha=10)
         la.fit(y_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: la.score(y_train,y_train)
```

```
Out[35]: 0.06110017645134225
```

## Accuracy(Lasso)

```
In [36]: la.score(y_test,y_test)
```

```
Out[36]: 0.060409984845479325
```

```
In [37]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(y_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef
```

```
Out[38]: array([-8.25695289e+00,  0.00000000e+00, -8.43038927e+00,  3.14630738e+00,
                4.20031356e-01,  0.00000000e+00,  1.95168770e-03,  3.63985159e+00,
               -1.15134089e-01,  3.03009505e-01,  2.87657806e+00, -4.44426535e-01,
                6.20824138e-01, -1.16271554e+00])
```

```
In [39]: en.intercept
```

```
Out[39]: 28079051.087767527
```

```
In [40]: prediction=en.predict(y_test)
```

```
In [41]: en.score(y_test,y_test)
```

```
Out[41]: 0.24396391772353465
```

## Evaluation Metrics

```
In [42]: from sklearn import metrics
         print(metrics.mean_absolute_error(y_test,prediction))
         print(metrics.mean_squared_error(y_test,prediction))
         print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))

32.078270309298276
1237.8837464625722
35.18357211061112
```

## Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
          'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df['station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (24758, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (24758,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[1,2,3,4,5,6,7,8,9,10,11,12,13,14]
```

```
In [51]: prediction=logr.predict(observation)  
print(prediction)  
[28079099]
```

```
In [52]: logr.classes
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8741416915744405
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 3.5557727473608076e-15
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[3.55577275e-15, 7.80743173e-29, 1.00000000e+00]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
```

```
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]  
                    }
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac  
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                'min_samples_leaf': [5, 10, 15, 20, 25],  
                                'n_estimators': [10, 20, 30, 40, 50]},  
                    scoring='accuracy')
```

```
In [60]: grid_search.best_score
```

```
Out[60]: 0.875072129255626
```

```
In [61]: rfc.best_estimator_
```



In [62]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
```

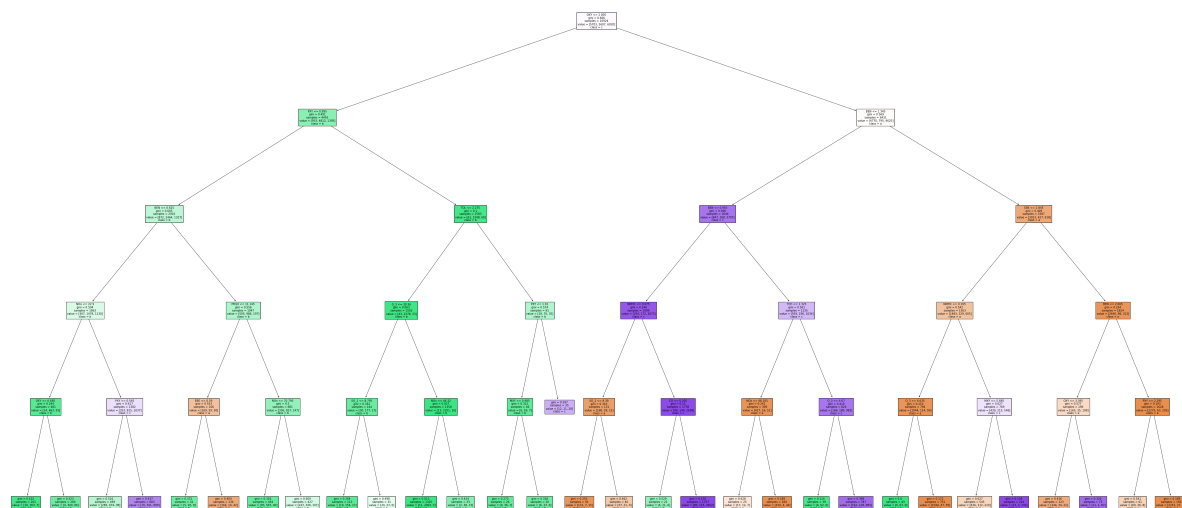
Out[62]: [Text(2222.7000000000003, 1993.2, 'OXY <= 1.005\ngini = 0.666\nsamples = 10924\nvalue = [5703, 5607, 6020]\nnclass = c'),  
Text(1171.8000000000002, 1630.8000000000002, 'PXY <= 0.995\ngini = 0.491\nsamples = 4493\nvalue = [933, 4812, 1395]\nnclass = b'),  
Text(595.2, 1268.4, 'BEN <= 0.625\ngini = 0.605\nsamples = 2910\nvalue = [872, 2464, 1327]\nnclass = b'),  
Text(297.6, 906.0, 'NOx <= 22.9\ngini = 0.594\nsamples = 1863\nvalue = [367, 1478, 1130]\nnclass = b'),  
Text(148.8, 543.5999999999999, 'OXY <= 0.685\ngini = 0.244\nsamples = 481\nvalue = [14, 663, 93]\nnclass = b'),  
Text(74.4, 181.19999999999982, 'gini = 0.102\nsamples = 201\nvalue = [10, 303, 7]\nnclass = b'),  
Text(223.20000000000002, 181.19999999999982, 'gini = 0.323\nsamples = 280\nvalue = [4, 360, 86]\nnclass = b'),  
Text(446.40000000000003, 543.5999999999999, 'PXY <= 0.545\ngini = 0.617\nsamples = 1382\nvalue = [353, 815, 1037]\nnclass = c'),  
Text(372.0, 181.19999999999982, 'gini = 0.516\nsamples = 499\nvalue = [283, 474, 38]\nnclass = b'),  
Text(520.80000000000001, 181.19999999999982, 'gini = 0.437\nsamples = 883\nvalue = [70, 341, 999]\nnclass = c'),  
Text(892.80000000000001, 906.0, 'PM10 <= 11.145\ngini = 0.556\nsamples = 1047\nvalue = [505, 986, 197]\nnclass = b'),  
Text(744.0, 543.5999999999999, 'EBE <= 0.59\ngini = 0.553\nsamples = 166\nvalue = [169, 59, 50]\nnclass = a'),  
Text(669.6, 181.19999999999982, 'gini = 0.372\nsamples = 32\nvalue = [5, 45, 8]\nnclass = b'),  
Text(818.40000000000001, 181.19999999999982, 'gini = 0.404\nsamples = 134\nvalue = [164, 14, 42]\nnclass = a'),  
Text(1041.60000000000001, 543.5999999999999, 'NOx <= 70.795\ngini = 0.5\nsamples = 881\nvalue = [336, 927, 147]\nnclass = b'),  
Text(967.2, 181.19999999999982, 'gini = 0.316\nsamples = 454\nvalue = [93, 587, 40]\nnclass = b'),  
Text(1116.0, 181.19999999999982, 'gini = 0.609\nsamples = 427\nvalue = [243, 340, 107]\nnclass = b'),  
Text(1748.4, 1268.4, 'TOL <= 2.275\ngini = 0.1\nsamples = 1583\nvalue = [61, 2348, 68]\nnclass = b'),  
Text(1488.0, 906.0, 'O\_3 <= 30.36\ngini = 0.063\nsamples = 1502\nvalue = [43, 2278, 33]\nnclass = b'),  
Text(1339.2, 543.5999999999999, 'SO\_2 <= 8.795\ngini = 0.352\nsamples = 144\nvalue = [30, 177, 17]\nnclass = b'),  
Text(1264.80000000000002, 181.19999999999982, 'gini = 0.264\nsamples = 113\nvalue = [10, 154, 17]\nnclass = b'),  
Text(1413.60000000000001, 181.19999999999982, 'gini = 0.498\nsamples = 31\nvalue = [20, 23, 0]\nnclass = b'),  
Text(1636.80000000000002, 543.5999999999999, 'NOx <= 44.32\ngini = 0.027\nsamples = 1358\nvalue = [13, 2101, 16]\nnclass = b'),  
Text(1562.4, 181.19999999999982, 'gini = 0.013\nsamples = 1325\nvalue = [11, 2063, 3]\nnclass = b'),  
Text(1711.2, 181.19999999999982, 'gini = 0.424\nsamples = 33\nvalue = [2, 38, 13]\nnclass = b'),  
Text(2008.80000000000002, 906.0, 'PXY <= 1.01\ngini = 0.574\nsamples = 81\nvalue = [10, 10, 10]\nnclass = c')]

```
lue = [18, 70, 35]\n\nclass = b'),  
Text(1934.4, 543.5999999999999, 'MXY <= 0.985\ngini = 0.312\nsamples = 46\nvalue = [6, 59, 7]\n\nclass = b'),  
Text(1860.0000000000002, 181.19999999999982, 'gini = 0.273\nsamples = 26\nvalue = [0, 36, 7]\n\nclass = b'),  
Text(2008.8000000000002, 181.19999999999982, 'gini = 0.328\nsamples = 20\nvalue = [6, 23, 0]\n\nclass = b'),  
Text(2083.2000000000003, 543.5999999999999, 'gini = 0.597\nsamples = 35\nvalue = [12, 11, 28]\n\nclass = c'),  
Text(3273.6000000000004, 1630.8000000000002, 'BEN <= 1.345\ngini = 0.569\nsamples = 6431\nvalue = [4770, 795, 4625]\n\nclass = a'),  
Text(2678.4, 1268.4, 'BEN <= 0.955\ngini = 0.398\nsamples = 3044\nvalue = [847, 368, 3707]\n\nclass = c'),  
Text(2380.8, 906.0, 'NMHC <= 0.075\ngini = 0.246\nsamples = 1909\nvalue = [254, 172, 2673]\n\nclass = c'),  
Text(2232.0, 543.5999999999999, 'SO_2 <= 8.39\ngini = 0.343\nsamples = 131\nvalue = [168, 28, 15]\n\nclass = a'),  
Text(2157.6000000000004, 181.19999999999982, 'gini = 0.255\nsamples = 91\nvalue = [131, 7, 15]\n\nclass = a'),  
Text(2306.4, 181.19999999999982, 'gini = 0.462\nsamples = 40\nvalue = [37, 21, 0]\n\nclass = a'),  
Text(2529.6000000000004, 543.5999999999999, 'CO <= 0.205\ngini = 0.15\nsamples = 1778\nvalue = [86, 144, 2658]\n\nclass = c'),  
Text(2455.2000000000003, 181.19999999999982, 'gini = 0.529\nsamples = 21\nvalue = [6, 21, 6]\n\nclass = b'),  
Text(2604.0, 181.19999999999982, 'gini = 0.135\nsamples = 1757\nvalue = [80, 123, 2652]\n\nclass = c'),  
Text(2976.0, 906.0, 'TCH <= 1.325\ngini = 0.561\nsamples = 1135\nvalue = [593, 196, 1034]\n\nclass = c'),  
Text(2827.2000000000003, 543.5999999999999, 'NOx <= 48.105\ngini = 0.241\nsamples = 309\nvalue = [427, 16, 51]\n\nclass = a'),  
Text(2752.8, 181.19999999999982, 'gini = 0.628\nsamples = 25\nvalue = [17, 12, 7]\n\nclass = a'),  
Text(2901.6000000000004, 181.19999999999982, 'gini = 0.189\nsamples = 284\nvalue = [410, 4, 44]\n\nclass = a'),  
Text(3124.8, 543.5999999999999, 'O_3 <= 4.67\ngini = 0.419\nsamples = 826\nvalue = [166, 180, 983]\n\nclass = c'),  
Text(3050.4, 181.19999999999982, 'gini = 0.114\nsamples = 39\nvalue = [4, 62, 0]\n\nclass = b'),  
Text(3199.2000000000003, 181.19999999999982, 'gini = 0.369\nsamples = 787\nvalue = [162, 118, 983]\n\nclass = c'),  
Text(3868.8, 1268.4, 'EBE <= 2.845\ngini = 0.409\nsamples = 3387\nvalue = [3923, 427, 918]\n\nclass = a'),  
Text(3571.2000000000003, 906.0, 'NMHC <= 0.195\ngini = 0.542\nsamples = 1563\nvalue = [1483, 329, 605]\n\nclass = a'),  
Text(3422.4, 543.5999999999999, 'O_3 <= 4.635\ngini = 0.253\nsamples = 794\nvalue = [1044, 114, 59]\n\nclass = a'),  
Text(3348.0000000000005, 181.19999999999982, 'gini = 0.0\nsamples = 43\nvalue = [0, 67, 0]\n\nclass = b'),  
Text(3496.8, 181.19999999999982, 'gini = 0.172\nsamples = 751\nvalue = [1044, 47, 59]\n\nclass = a'),  
Text(3720.0000000000005, 543.5999999999999, 'MXY <= 5.685\ngini = 0.627\nsamples = 769\nvalue = [439, 215, 546]\n\nclass = c'),  
Text(3645.6000000000004, 181.19999999999982, 'gini = 0.627\nsamples = 545\nvalue = [416, 212, 210]\n\nclass = a'),  
Text(3794.4, 181.19999999999982, 'gini = 0.134\nsamples = 224\nvalue = [23,
```

```

3, 336]\nclass = c'),
  Text(4166.400000000001, 906.0, 'BEN <= 2.065\ngini = 0.254\nsamples = 1824\nvalue = [2440, 98, 313]\nnclass = a'),
  Text(4017.6000000000004, 543.5999999999999, 'OXY <= 3.395\ngini = 0.577\nsamples = 196\nvalue = [165, 35, 108]\nnclass = a'),
  Text(3943.2000000000003, 181.19999999999982, 'gini = 0.436\nsamples = 123\nvalue = [144, 34, 21]\nnclass = a'),
  Text(4092.0000000000005, 181.19999999999982, 'gini = 0.326\nsamples = 73\nvalue = [21, 1, 87]\nnclass = c'),
  Text(4315.2000000000001, 543.5999999999999, 'PXY <= 2.295\ngini = 0.193\nsamples = 1628\nvalue = [2275, 63, 205]\nnclass = a'),
  Text(4240.8, 181.19999999999982, 'gini = 0.541\nsamples = 61\nvalue = [60, 36, 8]\nnclass = a'),
  Text(4389.6, 181.19999999999982, 'gini = 0.169\nsamples = 1567\nvalue = [221

```



## Conclusion

### Accuracy

**Linear Regression :0.38629705567566075**

**Ridge Regression : 0.06110017645134225**

**Lasso Regression : 0.060409984845479325**

**ElasticNet Regression : 0.24396391772353465**

**Logistic Regression : 0.8741416915744405**

**Random Forest :0.875072129255626**

**Random Forest is suitable for this dataset**