

# 20104016

## DEENA

### Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

### Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2007.csv")
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2007-12-01 01:00:00	NaN	2.86	NaN	NaN	NaN	282.200012	1054.000000	NaN	4.030000	15
1	2007-12-01 01:00:00	NaN	1.82	NaN	NaN	NaN	86.419998	354.600006	NaN	3.260000	8
2	2007-12-01 01:00:00	NaN	1.47	NaN	NaN	NaN	94.639999	319.000000	NaN	5.310000	5
3	2007-12-01 01:00:00	NaN	1.64	NaN	NaN	NaN	127.900002	476.700012	NaN	4.500000	10
4	2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	10
...	...	...	...	...	...	...	...	...	...	...	...
225115	2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	
225116	2007-03-01 00:00:00	NaN	0.16	NaN	NaN	NaN	46.820000	51.480000	NaN	22.150000	
225117	2007-03-01 00:00:00	0.24	NaN	0.20	NaN	0.09	51.259998	66.809998	NaN	18.540001	1
225118	2007-03-01 00:00:00	0.11	NaN	1.00	NaN	0.05	24.240000	36.930000	NaN	NaN	
225119	2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	1

225120 rows × 17 columns

# Data Cleaning and Data Preprocessing

In [3]: `df = df.dropna()`

In [4]: `df.columns`

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3',  
'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'],  
dtype='object')

In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        25443 non-null  object
1   BEN         25443 non-null  float64
2   CO          25443 non-null  float64
3   EBE         25443 non-null  float64
4   MXY         25443 non-null  float64
5   NMHC        25443 non-null  float64
6   NO_2        25443 non-null  float64
7   NOx         25443 non-null  float64
8   OXY         25443 non-null  float64
9   O_3         25443 non-null  float64
10  PM10        25443 non-null  float64
11  PM25        25443 non-null  float64
12  PXY         25443 non-null  float64
13  SO_2        25443 non-null  float64
14  TCH         25443 non-null  float64
15  TOL         25443 non-null  float64
16  station     25443 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

```
In [6]: data=df[['CO' , 'station']]
```

```
Out[6]:
```


	CO	station
4	1.86	28079006
21	0.31	28079024
25	1.42	28079099
30	1.89	28079006
47	0.30	28079024
...	...	...
225073	0.47	28079006
225094	0.45	28079099
225098	0.41	28079006
225115	0.45	28079024
225119	0.40	28079099

25443 rows × 2 columns

## Line chart

```
In [7]: data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:~>, <AxesSubplot:~>], dtype=object)
```



## Line chart

```
In [8]: data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```




## Bar chart

```
In [9]: k = data['CO:50']
```

```
In [10]: k.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```


```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: data.plot.area()
```


```
Out[12]: <AxesSubplot:>
```



## Box chart

In [13]: `data.plot.box()`


Out[13]: `<AxesSubplot:>`



## Pie chart

In [14]: `plot_stations()`


Out[14]: `<AxesSubplot:ylabel='station'>`



**Scatter chart**

In [15]: `data.plot.scatter(x='CO', y='station')`

Out[15]: `<AxesSubplot:xlabel='CO', ylabel='station'>`



In [16]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        25443 non-null  object
1   BEN         25443 non-null  float64
2   CO          25443 non-null  float64
3   EBE         25443 non-null  float64
4   MXY         25443 non-null  float64
5   NMHC        25443 non-null  float64
6   NO_2        25443 non-null  float64
7   NOx         25443 non-null  float64
8   OXY         25443 non-null  float64
9   O_3         25443 non-null  float64
10  PM10        25443 non-null  float64
11  PM25        25443 non-null  float64
12  PXY         25443 non-null  float64
13  SO_2        25443 non-null  float64
14  TCU         25443 non-null  float64
```



In [17]: `df.describe()`

Out[17]:

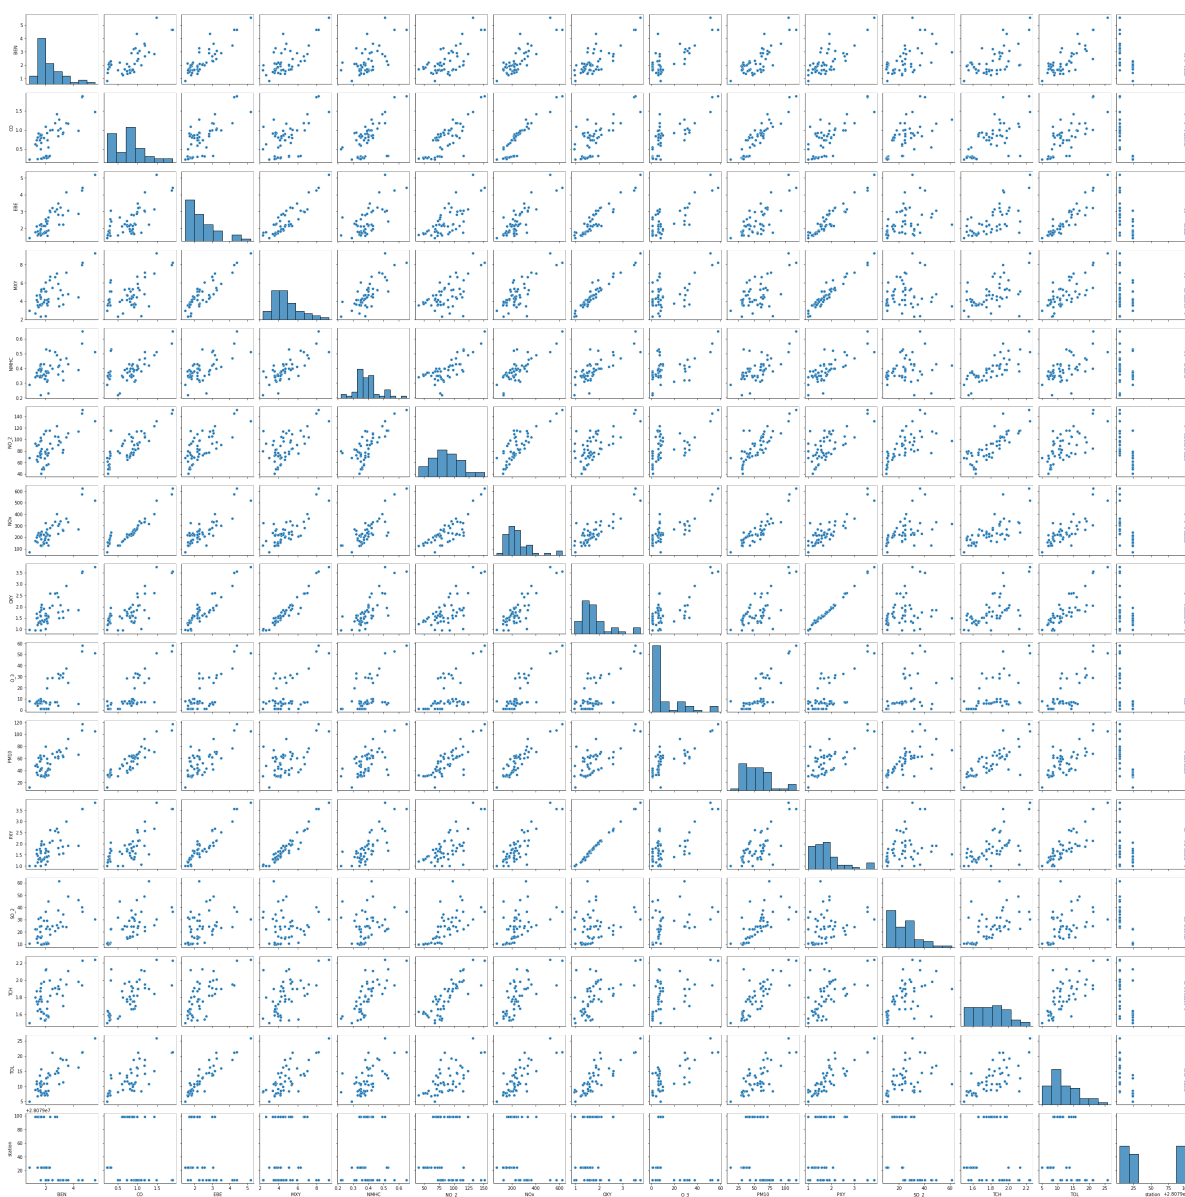
	BEN	CO	EBE	MXY	NMHC	NO_2	
<b>count</b>	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	254
<b>mean</b>	1.146744	0.505120	1.394071	2.392008	0.249967	58.532683	1
<b>std</b>	1.278733	0.423231	1.268265	2.784302	0.142627	37.755029	1
<b>min</b>	0.130000	0.000000	0.120000	0.150000	0.000000	1.690000	
<b>25%</b>	0.450000	0.260000	0.780000	0.960000	0.160000	31.285001	
<b>50%</b>	0.770000	0.400000	1.000000	1.500000	0.220000	54.080002	
<b>75%</b>	1.390000	0.640000	1.580000	2.855000	0.300000	79.230003	1
<b>max</b>	30.139999	9.660000	31.680000	65.480003	2.570000	430.299988	18

In [18]: `df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PM2.5', 'SO_2', 'TCH', 'TOL', 'Station']]`

## EDA AND VISUALIZATION

In [19]: `sns.pairplot(df1[0:50])`

Out[19]: `<seaborn.axisgrid.PairGrid at 0x1d69d281100>`




In [20]: `sns.distplot(df1['station'])`

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)

Out[20]: `<AxesSubplot:xlabel='station', ylabel='Density'>`



In [21]: `sns.heatmap(df1.corr())`

Out[21]: `<AxesSubplot:>`



## TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PM2_5', 'SO_2', 'TCH', 'TOL']]`  
`y=df['station']`

In [23]: `from sklearn.model_selection import train_test_split`  
`x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train, y_train)
```

Out[24]: LinearRegression()

```
In [25]: lr.intercept
```

Out[25]: 28079009.192362268

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```


Out[26]:

	Co-efficient
BEN	-32.438372
CO	18.343146
EBE	-0.101886
MXV	-1.087555
NMHC	-40.099581
NO_2	0.110310
NOx	-0.040358
OXY	5.547378
O_3	-0.016867
PM10	0.127035
PXY	7.050148
SO_2	0.136247
TCH	26.020745
TOL	3.252164

In [27]: `prediction = lr.predict(x_test)`

`plt.scatter(y_test, prediction)`

Out[27]: <matplotlib.collections.PathCollection at 0x1d6abc1cf70>



## ACCURACY

In [28]: `lr.score(y_test, y_test)`

Out[28]: 0.16379162671518677

In [29]: `lr.score(y_train, y_train)`

Out[29]: 0.15706639562597824

## Ridge and Lasso

In [30]: `from sklearn.linear_model import Ridge, Lasso`

In [31]: `rr=Ridge(alpha=10)`

`rr.fit(y_train, y_train)`

Out[31]: Ridge(alpha=10)

## Accuracy(Ridge)

In [32]: `rr.score(y_test, y_test)`

Out[32]: 0.16369555416612636

In [33]: `rr.score(y_train, y_train)`

Out[33]: 0.15701592078588478

```
In [34]: la=Lasso(alpha=10)
         la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.01247796461964934
```

## Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.014360120930396403
```

```
In [37]: from sklearn.linear_model import ElasticNet
         en=ElasticNet()
         en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: en.coef
```

```
Out[38]: array([-8.03719675,  0.          , -0.          ,  0.14618211, -0.          ,
                0.04709276, -0.04847005,  0.64915365, -0.0485963 ,  0.14996697,
                0.75777493, -0.01733057,  0.          ,  1.01190089])
```

```
In [39]: en.intercept
```

```
Out[39]: 28079045.901221745
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.06928545501091632
```

## Evaluation Metrics

```
In [42]: from sklearn import metrics
         print(metrics.mean_absolute_error(y_test,prediction))
         print(metrics.mean_squared_error(y_test,prediction))
         print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))

36.455485646200415
1521.148584724685
39.00190488584737
```

## Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
          'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
target_vector=df['Station']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (25443, 14)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (25443,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)  
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[1,2,3,4,5,6,7,8,9,10,11,12,13,14]
```

```
In [51]: prediction=logr.predict(observation)  
print(prediction)  
[28079099]
```

```
In [52]: logr.classes
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8146838030106512
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 1.082753977181323e-19
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[1.08275398e-19, 1.80383815e-19, 1.00000000e+00]])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
```

```
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [60]: grid_search.best_score
```

```
Out[60]: 0.8235261089275687
```

```
In [61]: rfc.best_estimator_
```



In [62]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
```

```
plot_tree(nfe_best_estimators_[1], feature_names=X.columns, class_names=['a', 'b'])
```

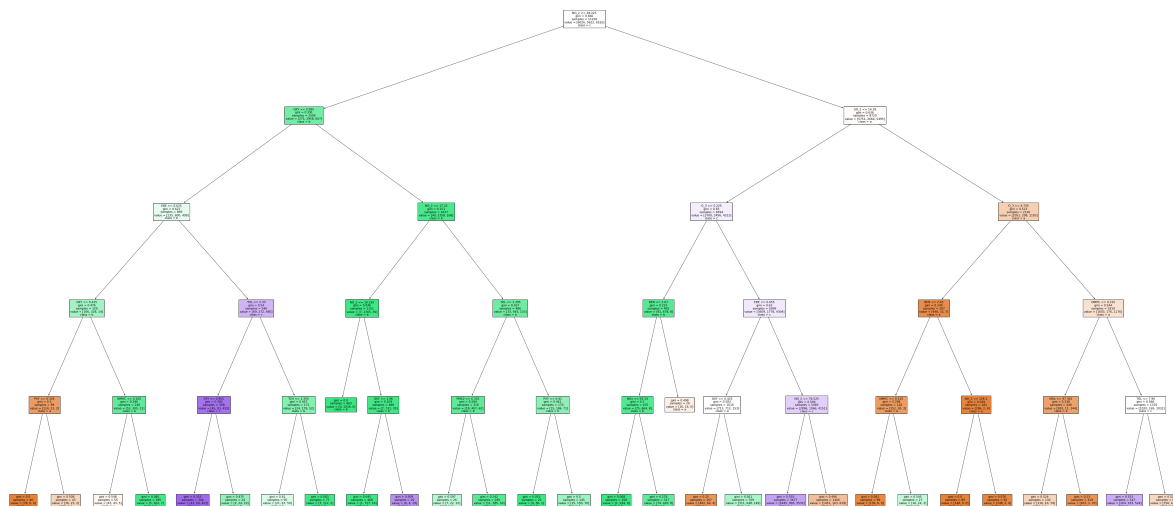
Out[62]: [Text(2192.142857142857, 1993.2, 'NO\_2 <= 28.225\ngini = 0.666\nsamples = 112  
26\nvalue = [6026, 5622, 6162]\nnclass = c'),  
Text(1135.9285714285713, 1630.8000000000002, 'OXY <= 0.995\ngini = 0.391\nsamples = 2506\nvalue = [275, 2958, 667]\nnclass = b'),  
Text(637.7142857142857, 1268.4, 'EBE <= 0.625\ngini = 0.627\nsamples = 869\nvalue = [235, 600, 499]\nnclass = b'),  
Text(318.85714285714283, 906.0, 'OXY <= 0.415\ngini = 0.476\nsamples = 329\nvalue = [166, 328, 14]\nnclass = b'),  
Text(159.42857142857142, 543.5999999999999, 'PXY <= 0.305\ngini = 0.3\nsamples = 89\nvalue = [114, 23, 2]\nnclass = a'),  
Text(79.71428571428571, 181.19999999999982, 'gini = 0.0\nsamples = 46\nvalue = [78, 0, 0]\nnclass = a'),  
Text(239.1428571428571, 181.19999999999982, 'gini = 0.508\nsamples = 43\nvalue = [36, 23, 2]\nnclass = a'),  
Text(478.2857142857142, 543.5999999999999, 'NMHC <= 0.165\ngini = 0.296\nsamples = 240\nvalue = [52, 305, 12]\nnclass = b'),  
Text(398.57142857142856, 181.19999999999982, 'gini = 0.548\nsamples = 55\nvalue = [47, 43, 5]\nnclass = a'),  
Text(558.0, 181.19999999999982, 'gini = 0.085\nsamples = 185\nvalue = [5, 26  
2, 7]\nnclass = b'),  
Text(956.5714285714284, 906.0, 'TOL <= 2.33\ngini = 0.54\nsamples = 540\nvalue = [69, 272, 485]\nnclass = c'),  
Text(797.1428571428571, 543.5999999999999, 'PXY <= 0.955\ngini = 0.392\nsamples = 368\nvalue = [45, 93, 433]\nnclass = c'),  
Text(717.4285714285713, 181.19999999999982, 'gini = 0.352\nsamples = 344\nvalue = [43, 69, 423]\nnclass = c'),  
Text(876.8571428571428, 181.19999999999982, 'gini = 0.475\nsamples = 24\nvalue = [2, 24, 10]\nnclass = b'),  
Text(1116.0, 543.5999999999999, 'TCH <= 1.355\ngini = 0.457\nsamples = 172\nvalue = [24, 179, 52]\nnclass = b'),  
Text(1036.2857142857142, 181.19999999999982, 'gini = 0.61\nsamples = 95\nvalue = [21, 67, 50]\nnclass = b'),  
Text(1195.7142857142856, 181.19999999999982, 'gini = 0.083\nsamples = 77\nvalue = [3, 112, 2]\nnclass = b'),  
Text(1634.142857142857, 1268.4, 'NO\_2 <= 17.22\ngini = 0.151\nsamples = 163  
7\nvalue = [40, 2358, 168]\nnclass = b'),  
Text(1355.142857142857, 906.0, 'NO\_2 <= 10.195\ngini = 0.046\nsamples = 115  
1\nvalue = [7, 1765, 35]\nnclass = b'),  
Text(1275.4285714285713, 543.5999999999999, 'gini = 0.0\nsamples = 663\nvalue = [0, 1034, 0]\nnclass = b'),  
Text(1434.8571428571427, 543.5999999999999, 'OXY <= 1.06\ngini = 0.104\nsamples = 488\nvalue = [7, 731, 35]\nnclass = b'),  
Text(1355.142857142857, 181.19999999999982, 'gini = 0.045\nsamples = 468\nvalue = [1, 727, 16]\nnclass = b'),  
Text(1514.5714285714284, 181.19999999999982, 'gini = 0.509\nsamples = 20\nvalue = [6, 4, 19]\nnclass = c'),  
Text(1913.1428571428569, 906.0, 'TOL <= 1.395\ngini = 0.357\nsamples = 486\nvalue = [33, 593, 133]\nnclass = b'),  
Text(1753.7142857142856, 543.5999999999999, 'PM10 <= 6.335\ngini = 0.284\nsamples = 316\nvalue = [18, 407, 62]\nnclass = b'),  
Text(1673.9999999999998, 181.19999999999982, 'gini = 0.597\nsamples = 26\nvalue = [1, 1, 1]\nnclass = a')]

```
lue = [7, 22, 12]\n\nclass = b'),
  Text(1833.4285714285713, 181.19999999999982, 'gini = 0.242\n\ nsamples = 290\n\ nvalue = [11, 385, 50]\n\nclass = b'),
  Text(2072.5714285714284, 543.59999999999999, 'PXY <= 0.41\n\ ngini = 0.461\n\ nsamples = 170\n\ nvalue = [15, 186, 71]\n\nclass = b'),
  Text(1992.8571428571427, 181.19999999999982, 'gini = 0.053\n\ nsamples = 25\n\ nvalue = [0, 36, 1]\n\nclass = b'),
  Text(2152.285714285714, 181.19999999999982, 'gini = 0.5\n\ nsamples = 145\n\ nvalue = [15, 150, 70]\n\nclass = b'),
  Text(3248.3571428571427, 1630.8000000000002, 'SO_2 <= 14.39\n\ ngini = 0.636\n\ nsamples = 8720\n\ nvalue = [5751, 2664, 5495]\n\nclass = a'),
  Text(2670.428571428571, 1268.4, 'O_3 <= 5.225\n\ ngini = 0.65\n\ nsamples = 6594\n\ nvalue = [3700, 2456, 4312]\n\nclass = c'),
  Text(2471.142857142857, 906.0, 'BEN <= 2.67\n\ ngini = 0.225\n\ nsamples = 495\n\ nvalue = [91, 678, 8]\n\nclass = b'),
  Text(2391.428571428571, 543.59999999999999, 'NOx <= 93.34\n\ ngini = 0.2\n\ nsamples = 475\n\ nvalue = [75, 664, 8]\n\nclass = b'),
  Text(2311.7142857142853, 181.19999999999982, 'gini = 0.008\n\ nsamples = 158\n\ nvalue = [1, 244, 0]\n\nclass = b'),
  Text(2471.142857142857, 181.19999999999982, 'gini = 0.278\n\ nsamples = 317\n\ nvalue = [74, 420, 8]\n\nclass = b'),
  Text(2550.8571428571427, 543.59999999999999, 'gini = 0.498\n\ nsamples = 20\n\ nvalue = [16, 14, 0]\n\nclass = a'),
  Text(2869.7142857142853, 906.0, 'EBE <= 0.655\n\ ngini = 0.63\n\ nsamples = 6099\n\ nvalue = [3609, 1778, 4304]\n\nclass = c'),
  Text(2710.285714285714, 543.59999999999999, 'OXY <= 0.425\n\ ngini = 0.583\n\ nsamples = 1016\n\ nvalue = [713, 712, 153]\n\nclass = a'),
  Text(2630.5714285714284, 181.19999999999982, 'gini = 0.25\n\ nsamples = 307\n\ nvalue = [402, 64, 4]\n\nclass = a'),
  Text(2790.0, 181.19999999999982, 'gini = 0.561\n\ nsamples = 709\n\ nvalue = [311, 648, 149]\n\nclass = b'),
  Text(3029.142857142857, 543.59999999999999, 'NO_2 <= 76.525\n\ ngini = 0.594\n\ nsamples = 5083\n\ nvalue = [2896, 1066, 4151]\n\nclass = c'),
  Text(2949.428571428571, 181.19999999999982, 'gini = 0.555\n\ nsamples = 3677\n\ nvalue = [1445, 903, 3532]\n\nclass = c'),
  Text(3108.8571428571427, 181.19999999999982, 'gini = 0.496\n\ nsamples = 1406\n\ nvalue = [1451, 163, 619]\n\nclass = a'),
  Text(3826.2857142857138, 1268.4, 'O_3 <= 4.705\n\ ngini = 0.523\n\ nsamples = 2126\n\ nvalue = [2051, 208, 1183]\n\nclass = a'),
  Text(3507.428571428571, 906.0, 'BEN <= 2.63\n\ ngini = 0.149\n\ nsamples = 307\n\ nvalue = [448, 32, 7]\n\nclass = a'),
  Text(3347.9999999999995, 543.59999999999999, 'NMHC <= 0.415\n\ ngini = 0.298\n\ nsamples = 126\n\ nvalue = [152, 30, 3]\n\nclass = a'),
  Text(3268.285714285714, 181.19999999999982, 'gini = 0.081\n\ nsamples = 99\n\ nvalue = [136, 6, 0]\n\nclass = a'),
  Text(3427.7142857142853, 181.19999999999982, 'gini = 0.545\n\ nsamples = 27\n\ nvalue = [16, 24, 3]\n\nclass = b'),
  Text(3666.8571428571427, 543.59999999999999, 'NO_2 <= 126.1\n\ ngini = 0.039\n\ nsamples = 181\n\ nvalue = [296, 2, 4]\n\nclass = a'),
  Text(3587.142857142857, 181.19999999999982, 'gini = 0.0\n\ nsamples = 89\n\ nvalue = [148, 0, 0]\n\nclass = a'),
  Text(3746.5714285714284, 181.19999999999982, 'gini = 0.076\n\ nsamples = 92\n\ nvalue = [148, 2, 4]\n\nclass = a'),
  Text(4145.142857142857, 906.0, 'NMHC <= 0.215\n\ ngini = 0.544\n\ nsamples = 1819\n\ nvalue = [1603, 176, 1176]\n\nclass = a'),
  Text(3985.7142857142853, 543.59999999999999, 'NOx <= 97.365\n\ ngini = 0.338\n\ nsamples = 1016\n\ nvalue = [713, 712, 153]\n\nclass = a')
```

```

mples = 449\nvalue = [583, 11, 144]\nnclass = a'),
Text(3905.9999999999995, 181.19999999999982, 'gini = 0.524\nsamples = 130\nv
alue = [116, 10, 74]\nnclass = a'),
Text(4065.428571428571, 181.19999999999982, 'gini = 0.23\nsamples = 319\nval
ue = [467, 1, 70]\nnclass = a'),
Text(4304.571428571428, 543.5999999999999, 'TOL <= 7.96\ngini = 0.566\nsampl
es = 1370\nvalue = [1020, 165, 1032]\nnclass = c'),
Text(4224.857142857142, 181.19999999999982, 'gini = 0.553\nsamples = 547\nva
lue = [264, 103, 524]\nnclass = c'),
Text(4384.285714285714, 181.19999999999982, 'gini = 0.526\nsamples = 823\nva

```



## Conclusion

### Accuracy

**Linear Regression :0.15706639562597824**

**Ridge Regression : 0.01247796461964934**

**Lasso Regression : 0.014360120930396403**

**ElasticNet Regression : 0.06928545501091632**

**Logistic Regression : 0.8146838030106512**

**Random Forest :0.8235261089275687**

**Random Forest is suitable for this dataset**

