

# 20104016

## DEENA

### Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
```

### Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2008.csv")
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.
...	...	...	...	...	...	...	...	...	...	...	...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.
226388	2008-11-01 00:00:00	NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN	35.830002	15.
226389	2008-11-01 00:00:00	0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN	14.130000	17.
226390	2008-11-01 00:00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN	NaN	11.
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.

226392 rows × 17 columns

# Data Cleaning and Data Preprocessing

In [3]:

In [4]:

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3',  
'PM10', 'PM25', 'PXY', 'SO\_2', 'TCH', 'TOL', 'station'],  
dtype='object')

In [5]:

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 25631 entries, 4 to 226391  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   date        25631 non-null  object  
1   BEN         25631 non-null  float64  
2   CO          25631 non-null  float64  
3   EBE         25631 non-null  float64  
4   MXY         25631 non-null  float64  
5   NMHC        25631 non-null  float64  
6   NO_2        25631 non-null  float64  
7   NOx         25631 non-null  float64  
8   OXY         25631 non-null  float64  
9   O_3         25631 non-null  float64  
10  PM10        25631 non-null  float64  
11  PM25        25631 non-null  float64  
12  PXY         25631 non-null  float64  
13  SO_2        25631 non-null  float64  
14  TCH         25631 non-null  float64  
15  TOL         25631 non-null  float64  
16  station     25631 non-null  int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 3.5+ MB
```

```
In [6]: data=df[['CO' , 'station']]
```

```
Out[6]:
```


	CO	station
4	0.80	28079006
21	0.37	28079024
25	0.39	28079099
30	0.51	28079006
47	0.39	28079024
...	...	...
226362	0.35	28079024
226366	0.46	28079099
226371	0.53	28079006
226387	0.30	28079024
226391	0.36	28079099

25631 rows × 2 columns

## Line chart

```
In [7]:
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

In [8]:

Out[8]: &lt;AxesSubplot:&gt;




## Bar chart

In [9]:

In [10]:


Out[10]: &lt;AxesSubplot:&gt;



## Histogram

In [11]:


Out[11]: <AxesSubplot:ylabel='Frequency'>



## Area chart

In [12]:


Out[12]: <AxesSubplot:>



## Box chart

In [13]:


Out[13]: &lt;AxesSubplot:&gt;



## Pie chart

In [14]:


Out[14]: <AxesSubplot:ylabel='station'>



**Scatter chart**

In [15]:

Out[15]: &lt;AxesSubplot:xlabel='CO', ylabel='station'&gt;



In [16]:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25631 entries, 4 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        25631 non-null  object
1   BEN         25631 non-null  float64
2   CO          25631 non-null  float64
3   EBE         25631 non-null  float64
4   MXY         25631 non-null  float64
5   NMHC        25631 non-null  float64
6   NO_2        25631 non-null  float64
7   NOx         25631 non-null  float64
8   OXY         25631 non-null  float64
9   O_3         25631 non-null  float64
10  PM10        25631 non-null  float64
11  PM25        25631 non-null  float64
12  PXY         25631 non-null  float64
13  SO_2        25631 non-null  float64
14  TCU         25631 non-null  float64
```



	BEN	CO	EBE	MXV	NMHC	NO_2	
count	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000
mean	1.090541	0.440632	1.352355	2.446045	0.213323	54.225261	
std	1.146461	0.317853	1.118191	2.390023	0.123409	38.164647	1.146461
min	0.100000	0.060000	0.170000	0.240000	0.000000	0.240000	
25%	0.430000	0.260000	0.740000	1.000000	0.130000	25.719999	
50%	0.750000	0.350000	1.000000	1.620000	0.190000	48.000000	
75%	1.320000	0.510000	1.580000	3.105000	0.270000	74.924999	1.320000
max	27.230000	7.030000	26.740000	55.889999	1.760000	554.900024	27.230000


```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
```

## EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df)
```


Out[19]: &lt;seaborn.axisgrid.PairGrid at 0x29e417b2be0&gt;



In [20]:


```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[20]: &lt;AxesSubplot:xlabel='station', ylabel='Density'&gt;



In [21]:

Out[21]: &lt;AxesSubplot:&gt;



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM2_5', 'SO_2', 'TCH', 'TOL']]
```

```
In [23]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
Out[24]: LinearRegression()
```

```
In [25]:
```

```
Out[25]: 28079031.439269386
```


```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[26]:
```

	Co-efficient
<b>BEN</b>	-25.935634
<b>CO</b>	-2.878921
<b>EBE</b>	-0.601321
<b>MXY</b>	7.361963
<b>NMHC</b>	-24.981870
<b>NO_2</b>	-0.007740
<b>NOx</b>	0.118970
<b>OXY</b>	3.609772
<b>O_3</b>	-0.122802
<b>PM10</b>	0.141612
<b>PXY</b>	2.563398
<b>SO_2</b>	-0.555249
<b>TCH</b>	18.891868
<b>TOL</b>	-1.891016

In [27]: `prediction = lr.predict(x_test)`

Out[27]: `<matplotlib.collections.PathCollection at 0x29e50178700>`



## ACCURACY

In [28]: `accuracy = accuracy_score(y_test, prediction)`

Out[28]: `0.14589220039799633`

In [29]: `accuracy = accuracy_score(y_test, prediction)`

Out[29]: `0.14249828065737014`

## Ridge and Lasso

In [30]: `from sklearn.linear_model import Ridge, Lasso`

In [31]: `rr=Ridge(alpha=10)`

Out[31]: `Ridge(alpha=10)`

## Accuracy(Ridge)

In [32]: `accuracy = accuracy_score(y_test, rr.predict(x_test))`

Out[32]: `0.14580956282768287`

In [33]: `accuracy = accuracy_score(y_test, rr.predict(x_test))`

Out[33]: `0.14247468266116436`

```
In [34]: la=Lasso(alpha=10)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]:
```

```
Out[35]: 0.04072393966873966
```

## Accuracy(Lasso)

```
In [36]:
```

```
Out[36]: 0.04357378194132566
```

```
In [37]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
Out[37]: ElasticNet()
```

```
In [38]:
```

```
Out[38]: array([-4.64736325, -0.          ,  0.          ,  3.28696488, -0.          ,
                0.077272   ,  0.01721591,  1.48519662, -0.14541509,  0.14244237,
                1.67756072, -0.91206623,  0.          , -2.5707591 ])
```

```
In [39]:
```

```
Out[39]: 28079056.107349645
```

```
In [40]:
```

```
In [41]:
```

```
Out[41]: 0.096440753633108
```

## Evaluation Metrics

```
In [42]: from sklearn import metrics
          print(metrics.mean_absolute_error(y_test,prediction))
          print(metrics.mean_squared_error(y_test,prediction))
```

```
35.69170876745575
```

```
1483.646482465846
```

```
38.51813186625029
```

## Logistic Regression

In [43]:

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
          'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
```

In [45]:

Out[45]: (25631, 14)

In [46]:

Out[46]: (25631,)

In [47]:

In [48]:

```
In [49]: logr=LogisticRegression(max_iter=10000)
```

Out[49]: LogisticRegression(max\_iter=10000)

In [50]:

```
In [51]: prediction=logr.predict(observation)  
[28079099]
```

In [52]:

Out[52]: array([28079006, 28079024, 28079099], dtype=int64)

In [53]:

Out[53]: 0.794194530061254

In [54]:

Out[54]: 8.321803242555043e-09

In [55]:

Out[55]: array([[8.32180324e-09, 1.19114634e-13, 9.99999992e-01]])

## Random Forest

In [56]:

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)  
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                'min_samples_leaf': [5, 10, 15, 20, 25],  
                                'n_estimators': [10, 20, 30, 40, 50]},  
                    scoring='accuracy')
```

```
In [60]:
```

```
Out[60]: 0.8505095074715543
```

```
In [61]:
```



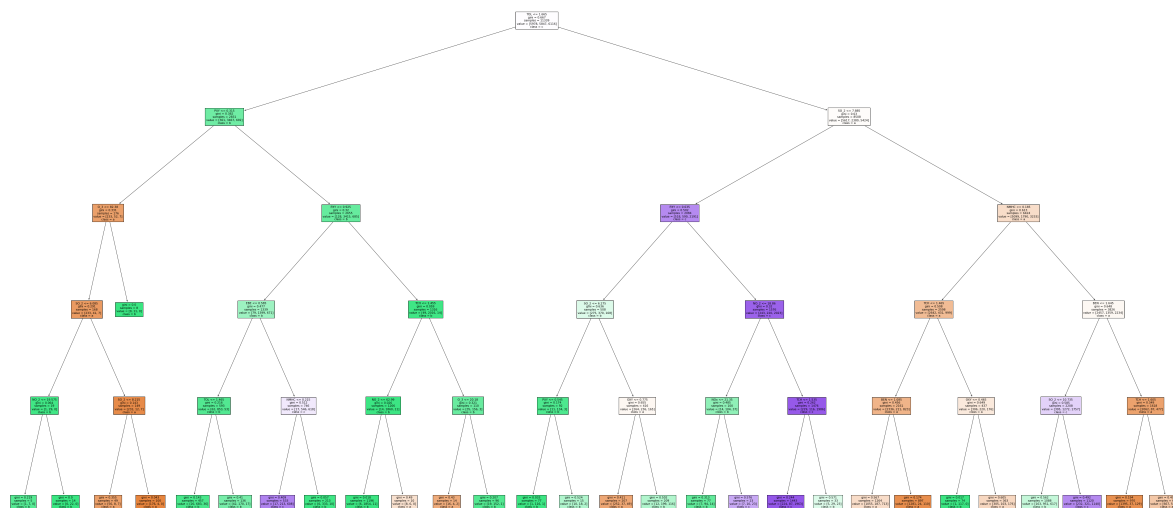
In [62]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
```

Out[62]: [Text(2012.7857142857142, 1993.2, 'TOL <= 1.665\ngini = 0.667\nsamples = 1133  
9\nvalue = [5978, 5847, 6116]\nnclass = c'),  
Text(836.9999999999999, 1630.8000000000002, 'PXY <= 0.315\ngini = 0.382\nsam  
ples = 2831\nvalue = [361, 3467, 692]\nnclass = b'),  
Text(398.57142857142856, 1268.4, 'O\_3 <= 82.38\ngini = 0.331\nsamples = 176\  
nvalue = [233, 52, 7]\nnclass = a'),  
Text(318.85714285714283, 906.0, 'SO\_2 <= 6.085\ngini = 0.291\nsamples = 168\  
nvalue = [233, 41, 7]\nnclass = a'),  
Text(159.42857142857142, 543.5999999999999, 'NO\_2 <= 19.575\ngini = 0.064\ns  
amples = 19\nvalue = [1, 29, 0]\nnclass = b'),  
Text(79.71428571428571, 181.19999999999982, 'gini = 0.219\nsamples = 5\nvalu  
e = [1, 7, 0]\nnclass = b'),  
Text(239.1428571428571, 181.19999999999982, 'gini = 0.0\nsamples = 14\nvalue  
= [0, 22, 0]\nnclass = b'),  
Text(478.2857142857142, 543.5999999999999, 'SO\_2 <= 8.215\ngini = 0.143\nsam  
ples = 149\nvalue = [232, 12, 7]\nnclass = a'),  
Text(398.57142857142856, 181.19999999999982, 'gini = 0.355\nsamples = 49\nva  
lue = [56, 8, 7]\nnclass = a'),  
Text(558.0, 181.19999999999982, 'gini = 0.043\nsamples = 100\nvalue = [176,  
4, 0]\nnclass = a'),  
Text(478.2857142857142, 906.0, 'gini = 0.0\nsamples = 8\nvalue = [0, 11, 0]\n  
nclass = b'),  
Text(1275.4285714285713, 1268.4, 'PXY <= 0.925\ngini = 0.32\nsamples = 2655\  
nvalue = [128, 3415, 685]\nnclass = b'),  
Text(956.5714285714284, 906.0, 'EBE <= 0.585\ngini = 0.477\nsamples = 1339\n  
value = [79, 1399, 671]\nnclass = b'),  
Text(797.1428571428571, 543.5999999999999, 'TOL <= 1.465\ngini = 0.216\nsamp  
les = 593\nvalue = [62, 853, 53]\nnclass = b'),  
Text(717.4285714285713, 181.19999999999982, 'gini = 0.143\nsamples = 457\nva  
lue = [20, 683, 36]\nnclass = b'),  
Text(876.8571428571428, 181.19999999999982, 'gini = 0.41\nsamples = 136\nval  
ue = [42, 170, 17]\nnclass = b'),  
Text(1116.0, 543.5999999999999, 'NMHC <= 0.215\ngini = 0.512\nsamples = 746\  
nvalue = [17, 546, 618]\nnclass = c'),  
Text(1036.2857142857142, 181.19999999999982, 'gini = 0.409\nsamples = 533\nv  
alue = [17, 213, 608]\nnclass = c'),  
Text(1195.7142857142856, 181.19999999999982, 'gini = 0.057\nsamples = 213\nv  
alue = [0, 333, 10]\nnclass = b'),  
Text(1594.2857142857142, 906.0, 'TCH <= 1.455\ngini = 0.059\nsamples = 1316\  
nvalue = [49, 2016, 14]\nnclass = b'),  
Text(1434.8571428571427, 543.5999999999999, 'NO\_2 <= 62.99\ngini = 0.026\nsa  
mples = 1206\nvalue = [14, 1860, 11]\nnclass = b'),  
Text(1355.142857142857, 181.19999999999982, 'gini = 0.018\nsamples = 1196\nv  
alue = [6, 1854, 11]\nnclass = b'),  
Text(1514.5714285714284, 181.19999999999982, 'gini = 0.49\nsamples = 10\nval  
ue = [8, 6, 0]\nnclass = a'),  
Text(1753.7142857142856, 543.5999999999999, 'O\_3 <= 20.18\ngini = 0.321\nsam  
ples = 110\nvalue = [35, 156, 3]\nnclass = b'),  
Text(1673.9999999999998, 181.19999999999982, 'gini = 0.43\nsamples = 14\nval  
ue = [16, 4, 2]\nnclass = a'),  
Text(1833.4285714285713, 181.19999999999982, 'gini = 0.207\nsamples = 96\nva  
lue = [19, 152, 1]\nnclass = b'),

```
Text(3188.5714285714284, 1630.8000000000002, 'SO_2 <= 7.885\ngini = 0.63\nsamples = 8508\nvalue = [5617, 2380, 5424]\nclass = a'),
Text(2550.8571428571427, 1268.4, 'PXY <= 0.635\ngini = 0.502\nsamples = 2084\nvalue = [518, 590, 2191]\nclass = c'),
Text(2232.0, 906.0, 'SO_2 <= 6.275\ngini = 0.636\nsamples = 508\nvalue = [275, 370, 168]\nclass = b'),
Text(2072.5714285714284, 543.5999999999999, 'PXY <= 0.545\ngini = 0.174\nsamples = 92\nvalue = [11, 134, 3]\nclass = b'),
Text(1992.8571428571427, 181.1999999999982, 'gini = 0.033\nsamples = 77\nvalue = [1, 116, 1]\nclass = b'),
Text(2152.285714285714, 181.1999999999982, 'gini = 0.524\nsamples = 15\nvalue = [10, 18, 2]\nclass = b'),
Text(2391.428571428571, 543.5999999999999, 'OXY <= 0.775\ngini = 0.655\nsamples = 416\nvalue = [264, 236, 165]\nclass = a'),
Text(2311.7142857142853, 181.1999999999982, 'gini = 0.411\nsamples = 207\nvalue = [252, 37, 49]\nclass = a'),
Text(2471.142857142857, 181.1999999999982, 'gini = 0.502\nsamples = 209\nvalue = [12, 199, 116]\nclass = b'),
Text(2869.7142857142853, 906.0, 'NO_2 <= 18.86\ngini = 0.32\nsamples = 1576\nvalue = [243, 220, 2023]\nclass = c'),
Text(2710.285714285714, 543.5999999999999, 'NOx <= 21.35\ngini = 0.485\nsamples = 100\nvalue = [14, 104, 37]\nclass = b'),
Text(2630.5714285714284, 181.1999999999982, 'gini = 0.313\nsamples = 77\nvalue = [7, 94, 14]\nclass = b'),
Text(2790.0, 181.1999999999982, 'gini = 0.576\nsamples = 23\nvalue = [7, 10, 23]\nclass = c'),
Text(3029.142857142857, 543.5999999999999, 'TCH <= 1.535\ngini = 0.262\nsamples = 1476\nvalue = [229, 116, 1986]\nclass = c'),
Text(2949.428571428571, 181.1999999999982, 'gini = 0.244\nsamples = 1443\nvalue = [224, 87, 1963]\nclass = c'),
Text(3108.8571428571427, 181.1999999999982, 'gini = 0.571\nsamples = 33\nvalue = [5, 29, 23]\nclass = b'),
Text(3826.2857142857138, 1268.4, 'NMHC <= 0.185\ngini = 0.613\nsamples = 6424\nvalue = [5099, 1790, 3233]\nclass = a'),
Text(3507.428571428571, 906.0, 'TCH <= 1.465\ngini = 0.508\nsamples = 2598\nvalue = [2642, 431, 999]\nclass = a'),
Text(3347.9999999999995, 543.5999999999999, 'BEN <= 1.085\ngini = 0.456\nsamples = 2161\nvalue = [2336, 211, 823]\nclass = a'),
Text(3268.285714285714, 181.1999999999982, 'gini = 0.567\nsamples = 1264\nvalue = [1055, 187, 713]\nclass = a'),
Text(3427.7142857142853, 181.1999999999982, 'gini = 0.174\nsamples = 897\nvalue = [1281, 24, 110]\nclass = a'),
Text(3666.8571428571427, 543.5999999999999, 'OXY <= 0.465\ngini = 0.649\nsamples = 437\nvalue = [306, 220, 176]\nclass = a'),
Text(3587.142857142857, 181.1999999999982, 'gini = 0.017\nsamples = 74\nvalue = [1, 117, 0]\nclass = b'),
Text(3746.5714285714284, 181.1999999999982, 'gini = 0.605\nsamples = 363\nvalue = [305, 103, 176]\nclass = a'),
Text(4145.142857142857, 906.0, 'BEN <= 1.645\ngini = 0.648\nsamples = 3826\nvalue = [2457, 1359, 2234]\nclass = a'),
Text(3985.7142857142853, 543.5999999999999, 'SO_2 <= 10.735\ngini = 0.585\nsamples = 2208\nvalue = [395, 1272, 1757]\nclass = c'),
Text(3905.9999999999995, 181.1999999999982, 'gini = 0.562\nsamples = 1088\nvalue = [163, 951, 617]\nclass = b'),
Text(4065.428571428571, 181.1999999999982, 'gini = 0.492\nsamples = 1120\nvalue = [232, 321, 1140]\nclass = c'),
```

```
Text(4304.571428571428, 543.5999999999999, 'TCH <= 1.605\ngini = 0.349\nsamples = 1618\nvalue = [2062, 87, 477]\nnclass = a'),
Text(4224.857142857142, 181.19999999999982, 'gini = 0.194\nsamples = 970\nvalue = [1395, 37, 129]\nnclass = a'),
Text(4384.285714285714, 181.19999999999982, 'gini = 0.499\nsamples = 648\nvalue = [667, 50, 348]\nnclass = a')]
```



## Conclusion

### Accuracy

**Linear Regression : 0.14249828065737014**

**Ridge Regression : 0.04072393966873966**

**Lasso Regression : 0.04357378194132566**

**ElasticNet Regression : 0.096440753633108**

**Logistic Regression : 0.794194530061254**

**Random Forest : 0.8505095074715543**

**Random Forest is suitable for this dataset**

