

20104016

DEENA

Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
```

Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2009.csv")
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.2
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.5
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.1
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.7
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.9
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.4
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3

215688 rows × 17 columns

Data Cleaning and Data Preprocessing

In [3]:

In [4]:

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
dtype='object')

In [5]:

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 24717 entries, 3 to 215687  
Data columns (total 17 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   date        24717 non-null  object  
1   BEN         24717 non-null  float64  
2   CO          24717 non-null  float64  
3   EBE         24717 non-null  float64  
4   MXY         24717 non-null  float64  
5   NMHC        24717 non-null  float64  
6   NO_2        24717 non-null  float64  
7   NOx         24717 non-null  float64  
8   OXY         24717 non-null  float64  
9   O_3         24717 non-null  float64  
10  PM10        24717 non-null  float64  
11  PM25        24717 non-null  float64  
12  PXY         24717 non-null  float64  
13  SO_2        24717 non-null  float64  
14  TCH         24717 non-null  float64  
15  TOL         24717 non-null  float64  
16  station     24717 non-null  int64  
dtypes: float64(15), int64(1), object(1)  
memory usage: 3.4+ MB
```

```
In [6]: data=df[['CO' , 'station']]
```

```
Out[6]:
```


	CO	station
3	0.33	28079006
20	0.32	28079024
24	0.24	28079099
28	0.21	28079006
45	0.30	28079024
...
215659	0.27	28079024
215663	0.35	28079099
215667	0.29	28079006
215683	0.22	28079024
215687	0.32	28079099

24717 rows × 2 columns

Line chart

```
In [7]:
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

In [8]:

Out[8]: <AxesSubplot:>




Bar chart

In [9]:

In [10]:


Out[10]: <AxesSubplot:>



Histogram

In [11]:


Out[11]: <AxesSubplot:ylabel='Frequency'>



Area chart

In [12]:


Out[12]: <AxesSubplot:>



Box chart

In [13]:


Out[13]: <AxesSubplot:>



Pie chart

In [14]:


Out[14]: <AxesSubplot:ylabel='station'>



Scatter chart

In [15]:

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



In [16]:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24717 entries, 3 to 215687
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        24717 non-null  object
1   BEN         24717 non-null  float64
2   CO          24717 non-null  float64
3   EBE         24717 non-null  float64
4   MXY         24717 non-null  float64
5   NMHC        24717 non-null  float64
6   NO_2        24717 non-null  float64
7   NOx         24717 non-null  float64
8   OXY         24717 non-null  float64
9   O_3         24717 non-null  float64
10  PM10        24717 non-null  float64
11  PM25        24717 non-null  float64
12  PXY         24717 non-null  float64
13  SO_2        24717 non-null  float64
14  TCU         24717 non-null  float64
```


	BEN	CO	EBE	MXV	NMHC	NO_2	
count	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000	24717.000000
mean	1.010583	0.448056	1.262430	2.244469	0.219582	55.563929	
std	1.007345	0.291706	1.074768	2.242214	0.141661	38.911677	
min	0.170000	0.060000	0.250000	0.240000	0.000000	0.600000	
25%	0.460000	0.270000	0.720000	0.990000	0.140000	26.510000	
50%	0.670000	0.370000	1.000000	1.490000	0.190000	47.930000	
75%	1.180000	0.570000	1.430000	2.820000	0.260000	76.269997	1
max	22.379999	5.570000	47.669998	56.500000	2.580000	477.399994	14

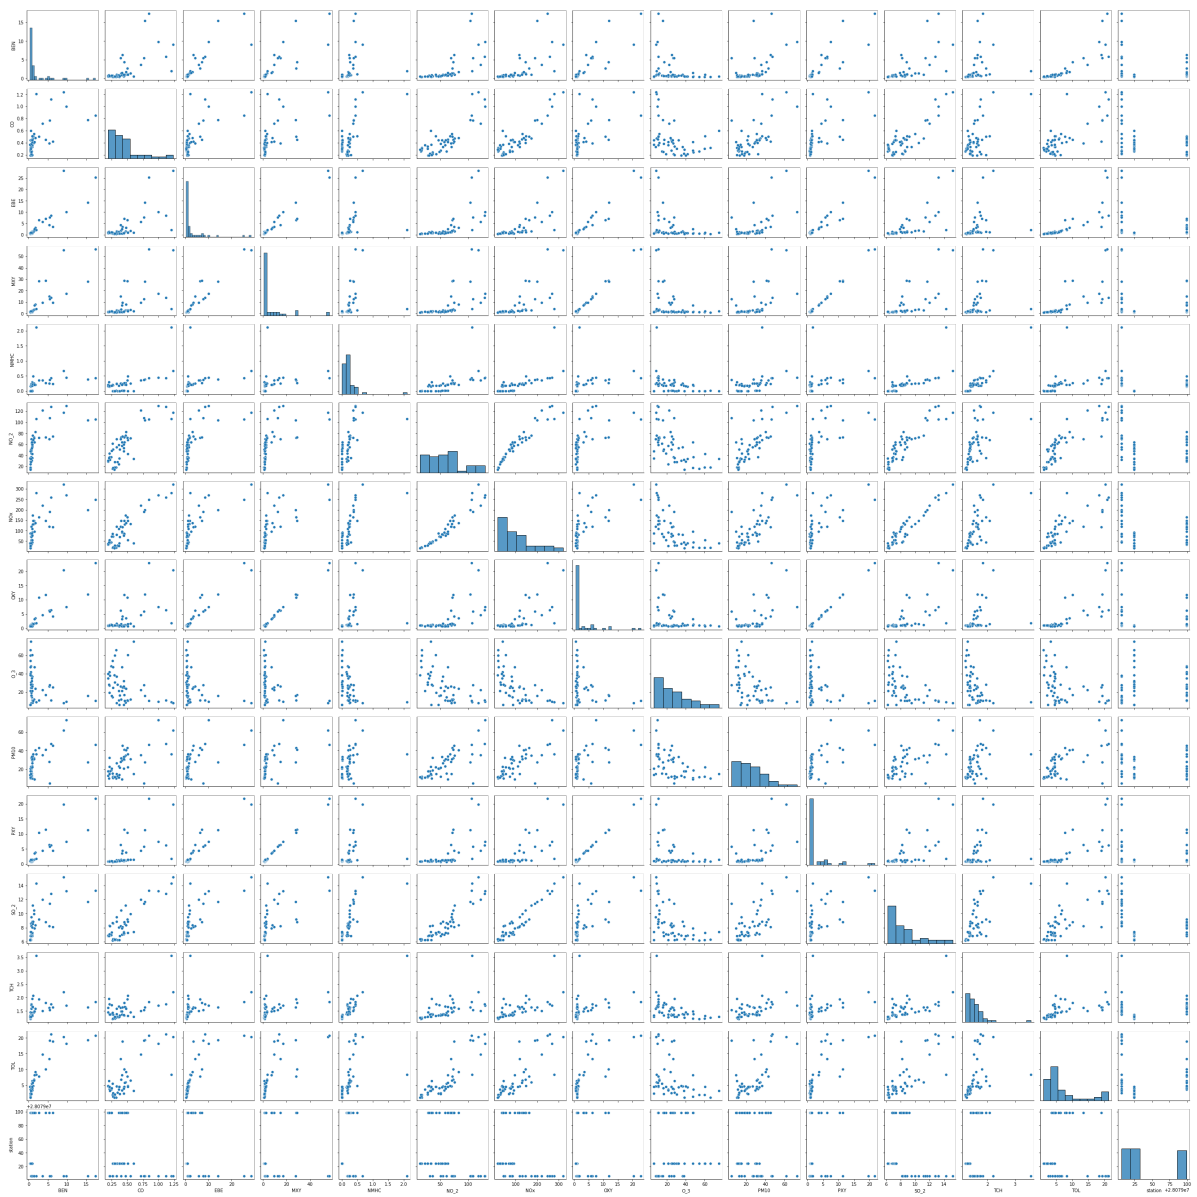
```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
       'PM10', 'PM25', 'SO_2', 'TSP', 'TSP1', 'TSP2', 'TSP3']]
```

EDA AND VISUALIZATION

In [19]:

`g = sns.pairplot(df)`


Out[19]: <seaborn.axisgrid.PairGrid at 0x19b08992280>



In [20]:


```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]:

Out[21]: <AxesSubplot:>



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM2.5', 'SO_2', 'TCH', 'TOL']]
```

```
In [23]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()
```

```
Out[24]: LinearRegression()
```

```
In [25]:
```

```
Out[25]: 28078910.857585404
```


```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[26]:
```

	Co-efficient
BEN	-36.203443
CO	-30.648171
EBE	6.488422
MXY	-0.905931
NMHC	-16.617460
NO_2	-0.170997
NOx	0.213958
OXY	13.118534
O_3	0.010307
PM10	-0.053681
PXY	3.870440
SO_2	-0.332929
TCH	111.370403
TOL	-1.288304

In [27]: `prediction = lr.predict(x_test)`

Out[27]: `<matplotlib.collections.PathCollection at 0x19b127da580>`



ACCURACY

In [28]:

Out[28]: `0.29852823000409645`

In [29]:

Out[29]: `0.2819204013446853`

Ridge and Lasso

In [30]:

In [31]: `rr=Ridge(alpha=10)`

Out[31]: `Ridge(alpha=10)`

Accuracy(Ridge)

In [32]:

Out[32]: `0.2967161015881199`

In [33]:

Out[33]: `0.2816595851212129`

```
In [34]: la=Lasso(alpha=10)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]:
```

```
Out[35]: 0.03737806206447625
```

Accuracy(Lasso)

```
In [36]:
```

```
Out[36]: 0.03407271577982274
```

```
In [37]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
Out[37]: ElasticNet()
```

```
In [38]:
```

```
Out[38]: array([-6.99357656, -0.70564139,  0.3459839 ,  2.25400319, -0.
               -0.20897486,  0.11926181,  1.21581674, -0.14525936,  0.07281443,
                1.9480131 , -0.77200717,  1.39621687, -2.06358417])
```

```
In [39]:
```

```
Out[39]: 28079063.86754315
```

```
In [40]:
```

```
In [41]:
```

```
Out[41]: 0.1051779453553251
```

Evaluation Metrics

```
In [42]: from sklearn import metrics
          print(metrics.mean_absolute_error(y_test,prediction))
          print(metrics.mean_squared_error(y_test,prediction))
```

```
35.91039339833247
```

```
1470.1634656890203
```

```
38.342710724321776
```

Logistic Regression

In [43]:

```
In [44]: feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O  
          'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
```

In [45]:

Out[45]: (24717, 14)

In [46]:

Out[46]: (24717,)

In [47]:

In [48]:

```
In [49]: logr=LogisticRegression(max_iter=10000)
```

Out[49]: LogisticRegression(max_iter=10000)

In [50]:

```
In [51]: prediction=logr.predict(observation)  
[28079099]
```

In [52]:

Out[52]: array([28079006, 28079024, 28079099], dtype=int64)

In [53]:

Out[53]: 0.8951733624630821

In [54]:

Out[54]: 5.447205522232353e-13

In [55]:

Out[55]: array([[5.44720552e-13, 8.28692830e-44, 1.00000000e+00]])

Random Forest

In [56]:

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)  
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],  
                    'min_samples_leaf':[5,10,15,20,25],  
                    'n_estimators':[10,20,30,40,50]
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
                    param_grid={'max_depth': [1, 2, 3, 4, 5],  
                                'min_samples_leaf': [5, 10, 15, 20, 25],  
                                'n_estimators': [10, 20, 30, 40, 50]},  
                    scoring='accuracy')
```

```
In [60]:
```

```
Out[60]: 0.8949199430985626
```

```
In [61]:
```


In [62]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
```

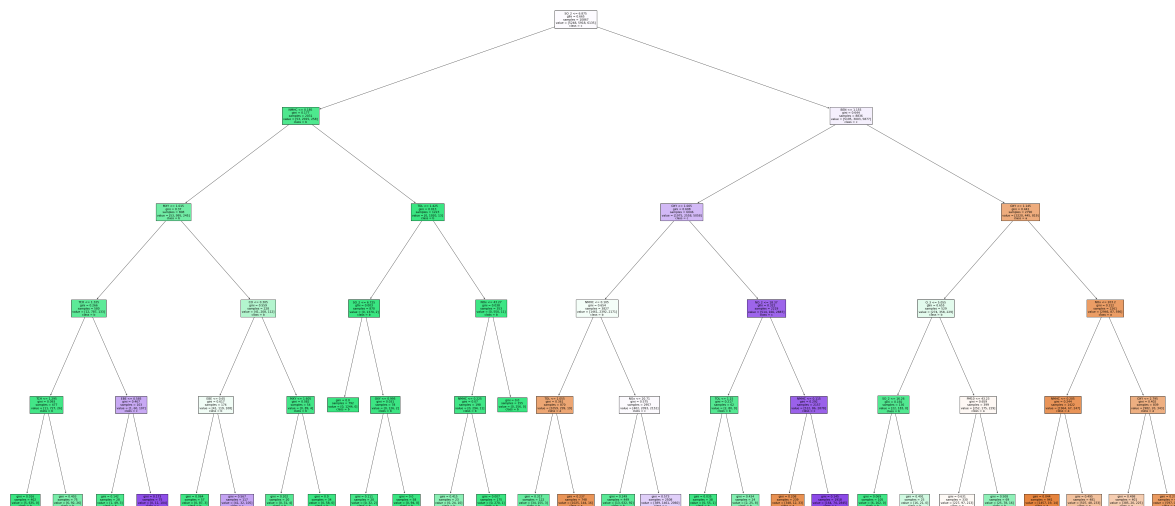
Out[62]: [Text(2152.285714285714, 1993.2, 'SO₂ <= 6.875\ngini = 0.665\nsamples = 1086
7\nvalue = [5248, 5918, 6135]\nnclass = c'),
Text(1116.0, 1630.8000000000002, 'NMHC <= 0.185\ngini = 0.177\nsamples = 203
1\nvalue = [53, 2915, 258]\nnclass = b'),
Text(637.7142857142857, 1268.4, 'MXY <= 1.015\ngini = 0.37\nsamples = 808\nvalue = [53, 995, 245]\nnclass = b'),
Text(318.85714285714283, 906.0, 'TCH <= 1.325\ngini = 0.266\nsamples = 580\nvalue = [12, 787, 133]\nnclass = b'),
Text(159.42857142857142, 543.5999999999999, 'TCH <= 1.295\ngini = 0.093\nsamples = 477\nvalue = [11, 727, 26]\nnclass = b'),
Text(79.71428571428571, 181.19999999999982, 'gini = 0.016\nsamples = 402\nvalue = [5, 635, 0]\nnclass = b'),
Text(239.1428571428571, 181.19999999999982, 'gini = 0.403\nsamples = 75\nvalue = [6, 92, 26]\nnclass = b'),
Text(478.2857142857142, 543.5999999999999, 'EBE <= 0.585\ngini = 0.467\nsamples = 103\nvalue = [1, 60, 107]\nnclass = c'),
Text(398.57142857142856, 181.19999999999982, 'gini = 0.142\nsamples = 28\nvalue = [1, 49, 3]\nnclass = b'),
Text(558.0, 181.19999999999982, 'gini = 0.173\nsamples = 75\nvalue = [0, 11, 104]\nnclass = c'),
Text(956.5714285714284, 906.0, 'CO <= 0.305\ngini = 0.559\nsamples = 228\nvalue = [41, 208, 112]\nnclass = b'),
Text(797.1428571428571, 543.5999999999999, 'EBE <= 0.65\ngini = 0.617\nsamples = 174\nvalue = [41, 119, 108]\nnclass = b'),
Text(717.4285714285713, 181.19999999999982, 'gini = 0.064\nsamples = 57\nvalue = [0, 87, 3]\nnclass = b'),
Text(876.8571428571428, 181.19999999999982, 'gini = 0.567\nsamples = 117\nvalue = [41, 32, 105]\nnclass = c'),
Text(1116.0, 543.5999999999999, 'MXY <= 1.605\ngini = 0.082\nsamples = 54\nvalue = [0, 89, 4]\nnclass = b'),
Text(1036.2857142857142, 181.19999999999982, 'gini = 0.202\nsamples = 20\nvalue = [0, 31, 4]\nnclass = b'),
Text(1195.7142857142856, 181.19999999999982, 'gini = 0.0\nsamples = 34\nvalue = [0, 58, 0]\nnclass = b'),
Text(1594.2857142857142, 1268.4, 'TOL <= 1.425\ngini = 0.013\nsamples = 122
3\nvalue = [0, 1920, 13]\nnclass = b'),
Text(1355.142857142857, 906.0, 'SO₂ <= 6.725\ngini = 0.003\nsamples = 870\nvalue = [0, 1370, 2]\nnclass = b'),
Text(1275.4285714285713, 543.5999999999999, 'gini = 0.0\nsamples = 792\nvalue = [0, 1244, 0]\nnclass = b'),
Text(1434.8571428571427, 543.5999999999999, 'OXY <= 0.995\ngini = 0.031\nsamples = 78\nvalue = [0, 126, 2]\nnclass = b'),
Text(1355.142857142857, 181.19999999999982, 'gini = 0.111\nsamples = 20\nvalue = [0, 32, 2]\nnclass = b'),
Text(1514.5714285714284, 181.19999999999982, 'gini = 0.0\nsamples = 58\nvalue = [0, 94, 0]\nnclass = b'),
Text(1833.4285714285713, 906.0, 'NOx <= 43.27\ngini = 0.038\nsamples = 353\nvalue = [0, 550, 11]\nnclass = b'),
Text(1753.7142857142856, 543.5999999999999, 'NMHC <= 0.225\ngini = 0.07\nsamples = 198\nvalue = [0, 294, 11]\nnclass = b'),
Text(1673.9999999999998, 181.19999999999982, 'gini = 0.415\nsamples = 23\nvalue = [0, 24, 10]\nnclass = b'),

```
Text(1833.4285714285713, 181.19999999999982, 'gini = 0.007\nsamples = 175\nvalue = [0, 270, 1]\nclass = b'),
Text(1913.1428571428569, 543.59999999999999, 'gini = 0.0\nsamples = 155\nvalue = [0, 256, 0]\nclass = b'),
Text(3188.5714285714284, 1630.80000000000002, 'BEN <= 1.155\ngini = 0.644\nsamples = 8836\nvalue = [5195, 3003, 5877]\nclass = c'),
Text(2550.8571428571427, 1268.4, 'OXY <= 1.005\ngini = 0.608\nsamples = 6046\nvalue = [1975, 2558, 5058]\nclass = c'),
Text(2232.0, 906.0, 'NMHC <= 0.105\ngini = 0.654\nsamples = 3827\nvalue = [1461, 2392, 2171]\nclass = b'),
Text(2072.5714285714284, 543.59999999999999, 'TOL <= 1.055\ngini = 0.361\nsamples = 870\nvalue = [1059, 299, 19]\nclass = a'),
Text(1992.8571428571427, 181.19999999999982, 'gini = 0.317\nsamples = 122\nvalue = [34, 155, 3]\nclass = b'),
Text(2152.285714285714, 181.19999999999982, 'gini = 0.237\nsamples = 748\nvalue = [1025, 144, 16]\nclass = a'),
Text(2391.428571428571, 543.59999999999999, 'NOx <= 20.71\ngini = 0.575\nsamples = 2957\nvalue = [402, 2093, 2152]\nclass = c'),
Text(2311.7142857142853, 181.19999999999982, 'gini = 0.249\nsamples = 449\nvalue = [13, 632, 92]\nclass = b'),
Text(2471.142857142857, 181.19999999999982, 'gini = 0.573\nsamples = 2508\nvalue = [389, 1461, 2060]\nclass = c'),
Text(2869.7142857142853, 906.0, 'NO_2 <= 18.37\ngini = 0.322\nsamples = 2219\nvalue = [514, 166, 2887]\nclass = c'),
Text(2710.285714285714, 543.59999999999999, 'TOL <= 1.15\ngini = 0.217\nsamples = 62\nvalue = [2, 80, 9]\nclass = b'),
Text(2630.5714285714284, 181.19999999999982, 'gini = 0.035\nsamples = 38\nvalue = [0, 55, 1]\nclass = b'),
Text(2790.0, 181.19999999999982, 'gini = 0.434\nsamples = 24\nvalue = [2, 25, 8]\nclass = b'),
Text(3029.142857142857, 543.59999999999999, 'NMHC <= 0.115\ngini = 0.292\nsamples = 2157\nvalue = [512, 86, 2878]\nclass = c'),
Text(2949.428571428571, 181.19999999999982, 'gini = 0.208\nsamples = 239\nvalue = [348, 12, 33]\nclass = a'),
Text(3108.8571428571427, 181.19999999999982, 'gini = 0.145\nsamples = 1918\nvalue = [164, 74, 2845]\nclass = c'),
Text(3826.2857142857138, 1268.4, 'OXY <= 1.145\ngini = 0.441\nsamples = 2790\nvalue = [3220, 445, 819]\nclass = a'),
Text(3507.428571428571, 906.0, 'O_3 <= 5.055\ngini = 0.655\nsamples = 529\nvalue = [274, 358, 229]\nclass = b'),
Text(3347.9999999999995, 543.59999999999999, 'SO_2 <= 16.26\ngini = 0.192\nsamples = 130\nvalue = [22, 183, 0]\nclass = b'),
Text(3268.285714285714, 181.19999999999982, 'gini = 0.069\nsamples = 105\nvalue = [6, 162, 0]\nclass = b'),
Text(3427.7142857142853, 181.19999999999982, 'gini = 0.491\nsamples = 25\nvalue = [16, 21, 0]\nclass = b'),
Text(3666.8571428571427, 543.59999999999999, 'PM10 <= 43.23\ngini = 0.659\nsamples = 399\nvalue = [252, 175, 229]\nclass = a'),
Text(3587.142857142857, 181.19999999999982, 'gini = 0.631\nsamples = 330\nvalue = [227, 97, 213]\nclass = a'),
Text(3746.5714285714284, 181.19999999999982, 'gini = 0.508\nsamples = 69\nvalue = [25, 78, 16]\nclass = b'),
Text(4145.142857142857, 906.0, 'NOx <= 203.2\ngini = 0.312\nsamples = 2261\nvalue = [2946, 87, 590]\nclass = a'),
Text(3985.7142857142853, 543.59999999999999, 'NMHC <= 0.205\ngini = 0.244\nsamples = 1422\nvalue = [1964, 67, 247]\nclass = a'),
```

```

Text(3905.9999999999995, 181.19999999999982, 'gini = 0.044\nsamples = 941\nvalue = [1457, 19, 14]\nnclass = a'),
Text(4065.428571428571, 181.19999999999982, 'gini = 0.495\nsamples = 481\nvalue = [507, 48, 233]\nnclass = a'),
Text(4304.571428571428, 543.5999999999999, 'OXY <= 2.795\ngini = 0.402\nsamples = 839\nvalue = [982, 20, 343]\nnclass = a'),
Text(4224.857142857142, 181.19999999999982, 'gini = 0.498\nsamples = 401\nvalue = [385, 20, 225]\nnclass = a'),
Text(4384.285714285714, 181.19999999999982, 'gini = 0.276\nsamples = 438\nvalue = [597, 0, 118]\nnclass = a')]

```



Conclusion

Accuracy

Linear Regression :0.2819204013446853

Ridge Regression : 0.03737806206447625

Lasso Regression : 0.03407271577982274

ElasticNet Regression :0.1051779453553251

Logistic Regression :0.8951733624630821

Random Forest :0.8949199430985626

Random Forest is suitable for this dataset

