

# 20104016

## DEENA

### Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
```

### Importing Datasets

```
In [2]: df=pd.read_csv("madrid_2014.csv")
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	3.0	10.0	NaN	NaN	NaN	3.0	NaN	NaN
1	2014-06-01 01:00:00	0.2	0.2	0.1	0.11	3.0	17.0	68.0	10.0	5.0	5.0	1.36	1.3
2	2014-06-01 01:00:00	0.3	NaN	0.1	NaN	2.0	6.0	NaN	NaN	NaN	NaN	NaN	1.1
3	2014-06-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	79.0	NaN	NaN	NaN	NaN	NaN
4	2014-06-01 01:00:00	NaN	NaN	NaN	NaN	1.0	6.0	75.0	NaN	NaN	4.0	NaN	NaN
...	...	...	...	...	...	...	...	...	...	...	...	...	...
210019	2014-09-01 00:00:00	NaN	0.5	NaN	NaN	20.0	84.0	29.0	NaN	NaN	NaN	NaN	NaN
210020	2014-09-01 00:00:00	NaN	0.3	NaN	NaN	1.0	22.0	NaN	15.0	NaN	6.0	NaN	NaN
210021	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	13.0	70.0	NaN	NaN	NaN	NaN	NaN
210022	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	38.0	42.0	NaN	NaN	NaN	NaN	NaN
210023	2014-09-01 00:00:00	NaN	NaN	NaN	NaN	1.0	26.0	65.0	11.0	NaN	NaN	NaN	NaN

210024 rows × 14 columns

## Data Cleaning and Data Preprocessing

In [3]:

In [4]:

Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO\_2', 'O\_3', 'PM10', 'PM25',  
'SO\_2', 'TCH', 'TOL', 'station'],  
dtype='object')

In [5]:

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 13946 entries, 1 to 210006  
Data columns (total 14 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   date        13946 non-null  object  
1   BEN         13946 non-null  float64  
2   CO          13946 non-null  float64  
3   EBE         13946 non-null  float64  
4   NMHC        13946 non-null  float64  
5   NO          13946 non-null  float64  
6   NO_2        13946 non-null  float64  
7   O_3         13946 non-null  float64  
8   PM10        13946 non-null  float64  
9   PM25        13946 non-null  float64  
10  SO_2        13946 non-null  float64  
11  TCH         13946 non-null  float64  
12  TOL         13946 non-null  float64  
13  station     13946 non-null  int64  
dtypes: float64(12), int64(1), object(1)  
memory usage: 1.6+ MB
```

```
In [6]: data=df[['CO' , 'station']]
```

```
Out[6]:
```


	CO	station
1	0.2	28079008
6	0.2	28079024
25	0.2	28079008
30	0.2	28079024
49	0.2	28079008
...	...	...
209958	0.2	28079024
209977	0.7	28079008
209982	0.2	28079024
210001	0.4	28079008
210006	0.2	28079024

13946 rows × 2 columns

## Line chart

```
In [7]:
```


```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

In [8]:

Out[8]: &lt;AxesSubplot:&gt;




## Bar chart

In [9]:

In [10]:


Out[10]: &lt;AxesSubplot:&gt;



## Histogram

In [11]:


Out[11]: &lt;AxesSubplot:ylabel='Frequency'&gt;



## Area chart

In [12]:


Out[12]: &lt;AxesSubplot:&gt;



## Box chart

In [13]:


Out[13]: &lt;AxesSubplot:&gt;



## Pie chart

In [14]:


Out[14]: <AxesSubplot:ylabel='station'>



**Scatter chart**

In [15]:

Out[15]: &lt;AxesSubplot:xlabel='CO', ylabel='station'&gt;



In [16]:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 13946 entries, 1 to 210006
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        13946 non-null  object
1   BEN         13946 non-null  float64
2   CO          13946 non-null  float64
3   EBE         13946 non-null  float64
4   NMHC        13946 non-null  float64
5   NO          13946 non-null  float64
6   NO_2        13946 non-null  float64
7   O_3         13946 non-null  float64
8   PM10        13946 non-null  float64
9   PM25        13946 non-null  float64
10  SO_2        13946 non-null  float64
11  TCH         13946 non-null  float64
12  TOL         13946 non-null  float64
13  station     13946 non-null  int64
```



	BEN	CO	EBE	NMHC	NO	NO_2
count	13946.000000	13946.000000	13946.000000	13946.000000	13946.000000	13946.000000
mean	0.375921	0.314793	0.306016	0.222302	17.589129	34.240929
std	0.555093	0.207375	0.635475	0.082403	39.432216	30.654229
min	0.100000	0.100000	0.100000	0.060000	1.000000	1.000000
25%	0.100000	0.200000	0.100000	0.160000	1.000000	10.000000
50%	0.200000	0.300000	0.100000	0.230000	4.000000	27.000000
75%	0.400000	0.400000	0.300000	0.260000	18.000000	51.000000
max	9.400000	4.400000	16.200001	1.290000	725.000000	346.000000


```
df1=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',  
        'PM10', 'SO_2', 'TSP', 'TSP1', 'TSP2', 'TSP3']]
```

## EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df)
```


Out[19]: &lt;seaborn.axisgrid.PairGrid at 0x2bee814c1f0&gt;



In [20]:


```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[20]: &lt;AxesSubplot:xlabel='station', ylabel='Density'&gt;



In [21]:

Out[21]: &lt;AxesSubplot:&gt;



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: x=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
               'PM10', 'SO_2', 'TCH', 'TOL']]
```

```
In [23]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

# Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression
         lr=LinearRegression()
```

Out[24]: LinearRegression()

```
In [25]: lr.intercept_
```

Out[25]: 28079022.114812426


```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

Out[26]:

	Co-efficient
<b>BEN</b>	-1.774907
<b>CO</b>	-5.818234
<b>EBE</b>	0.641452
<b>NMHC</b>	83.360064
<b>NO_2</b>	-0.032044
<b>O_3</b>	0.002743
<b>PM10</b>	0.015477
<b>SO_2</b>	-0.851834
<b>TCH</b>	-11.580837
<b>TOL</b>	-0.429355

```
In [27]: prediction =lr.predict(x_test)
```

Out[27]: <matplotlib.collections.PathCollection at 0x2bef1710460>



## ACCURACY

In [28]:

Out[28]: 0.8890761358656307

In [29]:

Out[29]: 0.8818792538063775

## Ridge and Lasso

In [30]:

In [31]: `rr=Ridge(alpha=10)`

Out[31]: Ridge(alpha=10)

## Accuracy(Ridge)

In [32]:

Out[32]: 0.8673565597079445

In [33]:

Out[33]: 0.8589350374750626

In [34]: `la=Lasso(alpha=10)`

Out[34]: Lasso(alpha=10)

In [35]:

Out[35]: 0.27689441343584587

## Accuracy(Lasso)

In [36]:

Out[36]: 0.2675964292464892

In [37]: `from sklearn.linear_model import ElasticNet  
en=ElasticNet()  
en.fit(x_train,y_train)`

Out[37]: ElasticNet()

In [38]:

```
Out[38]: array([ 0.          ,  0.          ,  0.16588122,  0.          , -0.04765468,
        -0.01300503,  0.02201435, -1.21102474,  0.          , -0.19119636])
```

In [39]:

```
Out[39]: 28079024.778188206
```

In [40]:

```
prediction=en.predict(x_test)
```

In [41]:

```
Out[41]: 0.47692368939175334
```

## Evaluation Metrics

In [42]:

```
from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
print(metrics.mean_squared_error(y_test,prediction))
```

```
5.022313464341247
```

```
33.01397887049697
```

```
5.745779222220166
```

## Logistic Regression

In [43]:

In [44]:

```
feature_matrix=df[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',
                  'PM10', 'SO_2', 'TCH', 'TOL']]
```

In [45]:

```
Out[45]: (13946, 10)
```

In [46]:

```
Out[46]: (13946,)
```

In [47]:

In [48]:

In [49]:

```
logr=LogisticRegression(max_iter=10000)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

In [50]:

In [51]:

```
prediction=logr.predict(observation)

[28079008]
```

In [52]:

Out[52]: array([28079008, 28079024], dtype=int64)

In [53]:

Out[53]: 0.9926143697117453

In [54]:

Out[54]: 1.0

In [55]:

Out[55]: array([[1.00000000e+00, 5.27113072e-18]])

## Random Forest

In [56]:

In [57]: rfc=RandomForestClassifier()

Out[57]: RandomForestClassifier()

In [58]:

```
parameters={'max_depth':[1,2,3,4,5],
            'min_samples_leaf':[5,10,15,20,25],
            'n_estimators':[10,20,30,40,50]}
```

In [59]:

```
from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="ac
```

Out[59]:

```
GridSearchCV(cv=2, estimator=RandomForestClassifier(),
             param_grid={'max_depth': [1, 2, 3, 4, 5],
                         'min_samples_leaf': [5, 10, 15, 20, 25],
                         'n_estimators': [10, 20, 30, 40, 50]},
             scoring='accuracy')
```

In [60]:

Out[60]: 0.9960049170251998

In [61]:

In [62]: `from sklearn.tree import plot_tree`

```
plt.figure(figsize=(80,40))
```

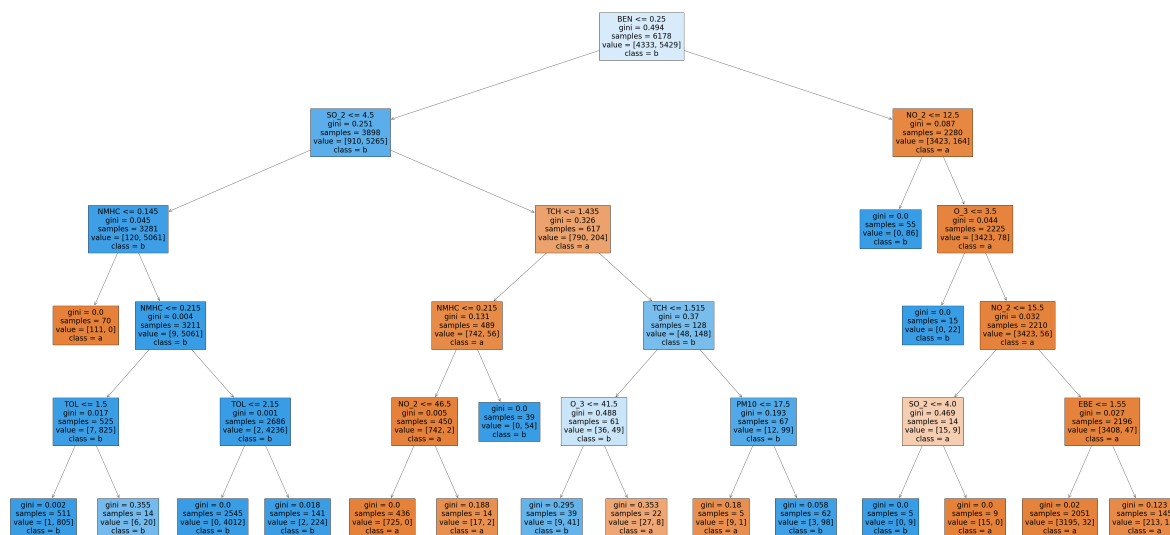
Out[62]: [Text(2411.3571428571427, 1993.2, 'BEN <= 0.25\ngini = 0.494\nsamples = 6178\nvalue = [4333, 5429]\nclass = b'),  
Text(1315.2857142857142, 1630.8000000000002, 'SO\_2 <= 4.5\ngini = 0.251\nsamples = 3898\nvalue = [910, 5265]\nclass = b'),  
Text(478.2857142857142, 1268.4, 'NMHC <= 0.145\ngini = 0.045\nsamples = 3281\nvalue = [120, 5061]\nclass = b'),  
Text(318.85714285714283, 906.0, 'gini = 0.0\nsamples = 70\nvalue = [111, 0]\nclass = a'),  
Text(637.7142857142857, 906.0, 'NMHC <= 0.215\ngini = 0.004\nsamples = 3211\nvalue = [9, 5061]\nclass = b'),  
Text(318.85714285714283, 543.5999999999999, 'TOL <= 1.5\ngini = 0.017\nsamples = 525\nvalue = [7, 825]\nclass = b'),  
Text(159.42857142857142, 181.19999999999982, 'gini = 0.002\nsamples = 511\nvalue = [1, 805]\nclass = b'),  
Text(478.2857142857142, 181.19999999999982, 'gini = 0.355\nsamples = 14\nvalue = [6, 20]\nclass = b'),  
Text(956.5714285714284, 543.5999999999999, 'TOL <= 2.15\ngini = 0.001\nsamples = 2686\nvalue = [2, 4236]\nclass = b'),  
Text(797.1428571428571, 181.19999999999982, 'gini = 0.0\nsamples = 2545\nvalue = [0, 4012]\nclass = b'),  
Text(1116.0, 181.19999999999982, 'gini = 0.018\nsamples = 141\nvalue = [2, 224]\nclass = b'),  
Text(2152.285714285714, 1268.4, 'TCH <= 1.435\ngini = 0.326\nsamples = 617\nvalue = [790, 204]\nclass = a'),  
Text(1753.7142857142856, 906.0, 'NMHC <= 0.215\ngini = 0.131\nsamples = 489\nvalue = [742, 56]\nclass = a'),  
Text(1594.2857142857142, 543.5999999999999, 'NO\_2 <= 46.5\ngini = 0.005\nsamples = 450\nvalue = [742, 2]\nclass = a'),  
Text(1434.8571428571427, 181.19999999999982, 'gini = 0.0\nsamples = 436\nvalue = [725, 0]\nclass = a'),  
Text(1753.7142857142856, 181.19999999999982, 'gini = 0.188\nsamples = 14\nvalue = [17, 2]\nclass = a'),  
Text(1913.1428571428569, 543.5999999999999, 'gini = 0.0\nsamples = 39\nvalue = [0, 54]\nclass = b'),  
Text(2550.8571428571427, 906.0, 'TCH <= 1.515\ngini = 0.37\nsamples = 128\nvalue = [48, 148]\nclass = b'),  
Text(2232.0, 543.5999999999999, 'O\_3 <= 41.5\ngini = 0.488\nsamples = 61\nvalue = [36, 49]\nclass = b'),  
Text(2072.5714285714284, 181.19999999999982, 'gini = 0.295\nsamples = 39\nvalue = [9, 41]\nclass = b'),  
Text(2391.428571428571, 181.19999999999982, 'gini = 0.353\nsamples = 22\nvalue = [27, 8]\nclass = a'),  
Text(2869.7142857142853, 543.5999999999999, 'PM10 <= 17.5\ngini = 0.193\nsamples = 67\nvalue = [12, 99]\nclass = b'),  
Text(2710.285714285714, 181.19999999999982, 'gini = 0.18\nsamples = 5\nvalue = [9, 1]\nclass = a'),  
Text(3029.142857142857, 181.19999999999982, 'gini = 0.058\nsamples = 62\nvalue = [3, 98]\nclass = b'),  
Text(3507.428571428571, 1630.8000000000002, 'NO\_2 <= 12.5\ngini = 0.087\nsamples = 2280\nvalue = [3423, 164]\nclass = a'),  
Text(3347.9999999999995, 1268.4, 'gini = 0.0\nsamples = 55\nvalue = [0, 86]\nclass = b'),



```

Text(3666.8571428571427, 1268.4, 'O_3 <= 3.5\ngini = 0.044\nsamples = 2225\n
value = [3423, 78]\nnclass = a'),
Text(3507.428571428571, 906.0, 'gini = 0.0\nsamples = 15\nvalue = [0, 22]\nc
lass = b'),
Text(3826.2857142857138, 906.0, 'NO_2 <= 15.5\ngini = 0.032\nsamples = 2210\
nvalue = [3423, 56]\nnclass = a'),
Text(3507.428571428571, 543.5999999999999, 'SO_2 <= 4.0\ngini = 0.469\nsampl
es = 14\nvalue = [15, 9]\nnclass = a'),
Text(3347.9999999999995, 181.19999999999982, 'gini = 0.0\nsamples = 5\nvalue
= [0, 9]\nnclass = b'),
Text(3666.8571428571427, 181.19999999999982, 'gini = 0.0\nsamples = 9\nvalue
= [15, 0]\nnclass = a'),
Text(4145.142857142857, 543.5999999999999, 'EBE <= 1.55\ngini = 0.027\nsampl
es = 2196\nvalue = [3408, 47]\nnclass = a'),
Text(3985.7142857142853, 181.19999999999982, 'gini = 0.02\nsamples = 2051\nv
alue = [3195, 32]\nnclass = a'),
Text(4304.571428571428, 181.19999999999982, 'gini = 0.123\nsamples = 145\nva
lue = [213, 15]\nnclass = a')]

```



## Conclusion

## Accuracy

**Linear Regression :0.8818792538063775**

**Ridge Regression :0.8673565597079445**

**Lasso Regression :0.2675964292464892**

**ElasticNet Regression : 0.47692368939175334**

**Logistic Regression : 0.9926143697117453**

Random Forest is 0.8888888888888888

**Random Forest is suitable for this dataset**